

1. Introduction

In this project, we are going to convert the sequential Heat Distribution program provided to parallel version with the help of MPI, Pthread, OpenMP, CUDA and MPI with OpenMP. This report will give the details of the design of each version and compare their performance.

2. Methods

In the given sequential file, the program would basically calculate and update the new temperature for each point in the grid using the information of its neighbors.

a) Conduct resource-consuming computing step in parallel

Throughout the program, the part requiring the most computing resources is the calculation part where the program would iterate through all the points in the grid and use the information of its neighbors to generate new temperature for it. Actually, this calculation is perfect for parallel computing since in one iteration, the calculations for each point are independent. Thus, we can simply just assign the task to every process or thread and finally gather the calculation result from them.

b) Divide the task into several packages

Down to the details, we need a way to split the task into packages with approximately equal size. Assume the size of the grid in grid is $size$ (that is, there are totally $size * size$ points in the grid). Assume that the total number of processes or threads in the program is $totalTaskSize$, we apply the following formula to divide the task. For process (thread) with task number i , it will calculate the point with from the start index to the end index (start and end index included):

$$StartIndex_i = \left(\frac{size * size}{totalTaskSize} \right) * i + min((size * size) \% totalTaskSize, i)$$

$$EndIndex_i = \left(\frac{size * size}{totalTaskSize} \right) * (i + 1) + min((size * size) \% totalTaskSize, i + 1) - 1$$

For instance, with $size$ equals to 20, there will be $20 * 20$ points in the grid, three processes would proceed the points at [0,133], [134,266], [267,399] respectively. In this case, the sizes of the sub-tasks are 133, 132 and 132 separately, which are approximately same.

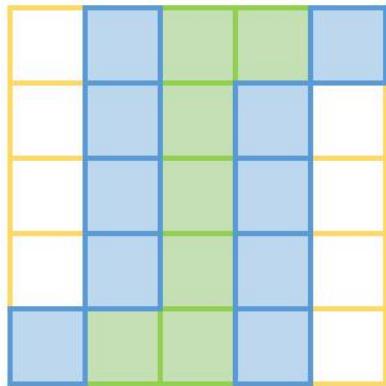
c) Mechanism of plotting with updated parameters in GUI

Notice that the calculation in the program is repeating endlessly since it always need to update the value of points after adjusting some special parameters (for example, `roo_size`) and replot the whole diagram. Thus, we need to make sure that other processes or threads would interact with the master process and thread correctly to compute with updated parameters.

d) Extra calculation of points in Sor

Different from Jacobi, Sor algorithm need to compute the points in the grid two rounds in each iteration and the second round need to use the information of the first round. This would be totally fine in multi-thread programming since after every thread reach the barrier at the end of the first round, all threads could easily use the calculated result of all points in the first round easily. However, in the multi-process programming, things become a little different since all process maintain their own local variables. To solve this problem, we could directly ask rank 0 to gather the

results from all processes and then broadcast all the calculation results to other processes though it will waste lot of time. Actually, an alternative way to solve this problem is to calculate the extra points in the first round to prepare for the second round. Basically, we just need to take care of the points around the calculating areas. One example to illustrate this is shown below:



Assumed the task assigned for this process is green squares, then it needs to calculate the blue squares as well in the first iteration in Sor to prepare for the second iteration.

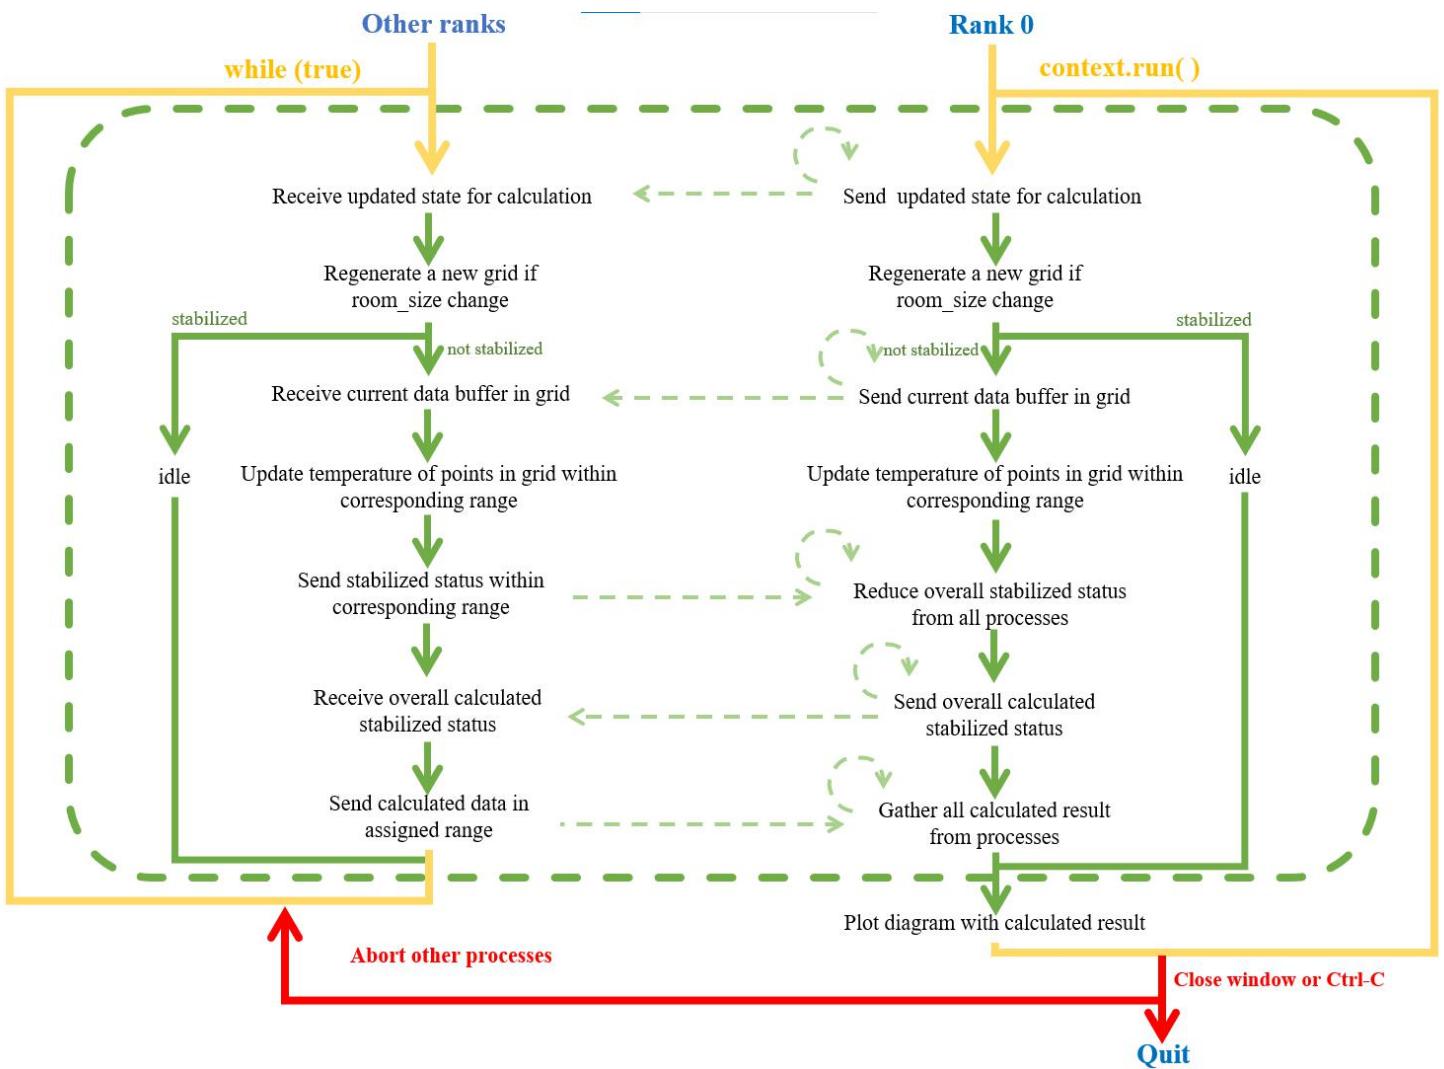
e) Solve data race in multi-thread programs

There is one important part may cause data race in multi-thread programs, which is to check all threads' status of stabilized. To solve this problem, the program applies a vector to store every thread's status and check all of them to generate the final result. However, if directly using `std::vector<bool>` it may still cause some extra data race since C++ adopt a special design for vector where only one bit is used to represent a bool for the sake of saving memory. Then, there may be the chance that multiple threads are writing one byte to change their bit allocated for them which is then the data race. To solve this problem, `std::vector<int>` instead is used as alternative.

d) MPI

Since we use rank 0 process as the master program which generate GUI and fetch the updated parameters from users, we need to make sure rank 0 pass necessary information to other ranks to do the corresponding calculation work. Following procedures are designed to make sure each process does the calculation correctly when doing calculation.

- Step 1: Rank 0 broadcast the state information of the current grid to each process.
- Step 2: All processes received the state information from rank 0. Then, all processes will check whether the room size has changed, if so, they will regenerate a new grid.
- Step 3: All process would calculate the next temperature within their corresponding range.
- Step 4: Rank 0 would reduce the stabilized information of each process local stabilized situation and generate the overall stabilized result which would then broadcast to all other processes.
- Step 5: Rank 0 would gather the calculated result from all other processes after which the Rank 0 would plot the heat grid. Next, all process would go to step 1.
- Step 6: If GUI window is closed or Ctrl-C is typed in the terminal, the rank 0 process would directly jump to step 7 and abort all other processes. In this case, you may see the abort signal prompted in the terminal which is a normal operation.



f) Pthread/OpenMP/CUDA

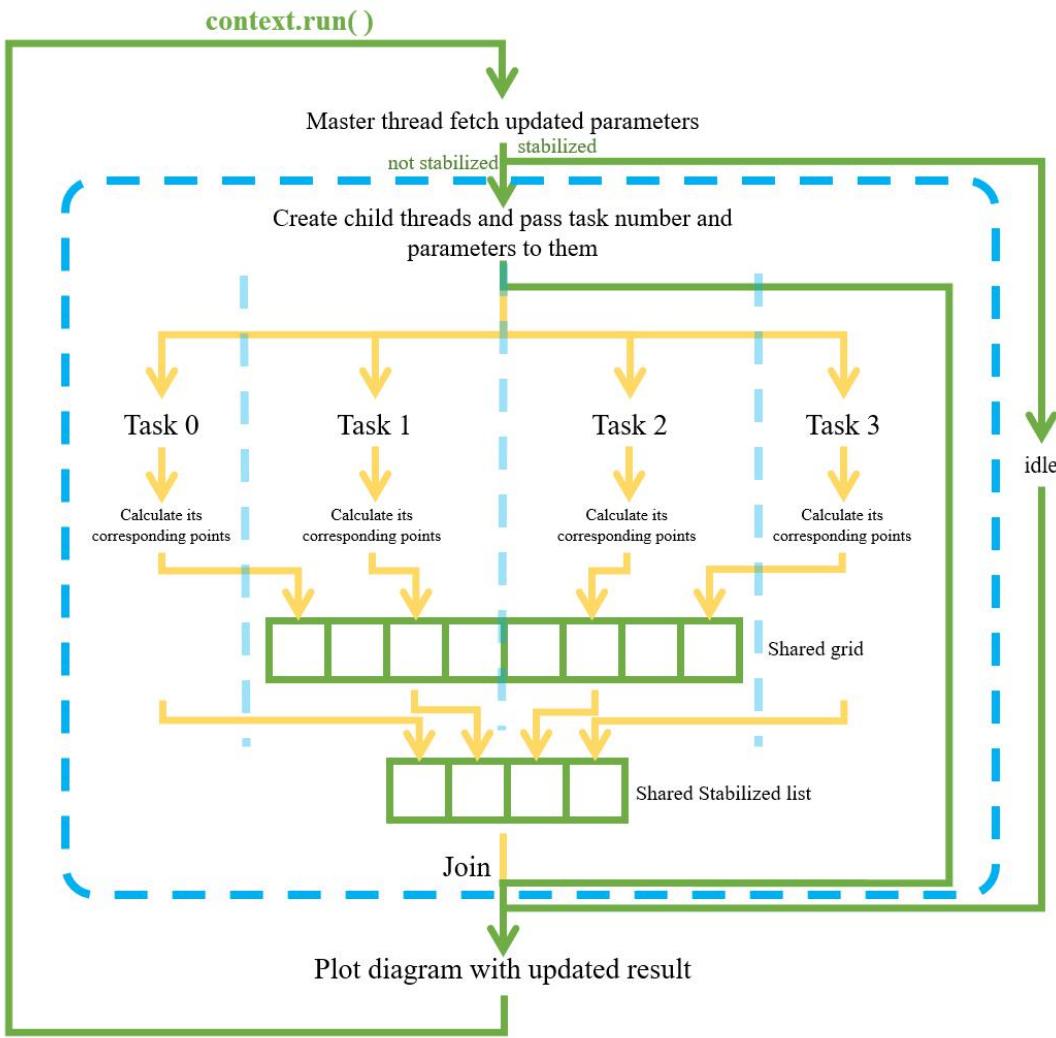
Besides using multi-processes, we can also use multi-threads to finish the calculation task. In this project, the Pthread, OpenMP and CUDA are implemented mainly follow the same logic.

With multiple threads, updating the diagram with new parameters would be much easier since we can directly pass parameters to threads when creating them. Waiting for all the threads finish their calculation task, the master thread would then plot the diagram directly. Followings are the main steps:

Step 1: Master program would create several threads and pass the corresponding parameters to them.

Step 2: All threads do the calculation next temperature on their corresponding regions. After that, they will write their local stabilized status into their corresponding positions on the status list.

Step 3: After joining all the threads, the master thread would then go through all the local status in the status list to generate the overall stabilized status. Finally, it will plot the heat grid and jump to step 1 again.



g) MPI with OpenMP

We can actually combine the multi-process and multi-thread mechanism together, that is, launch several processes and create multiple threads to finish the tasks assigned to each process.

Followings are the detailed steps which is just a simple combination of MPI and OpenMP:

Step 1: Rank 0 broadcast the state information of the current grid to each process.

Step 2: All processes received the state information from rank 0. Then, all processes will check whether the room size has changed, if so, they will regenerate a new grid.

Step 3: All process would calculate the next temperature within their corresponding range.

Step 4: Master program in each process would create several threads and pass the corresponding parameters to them.

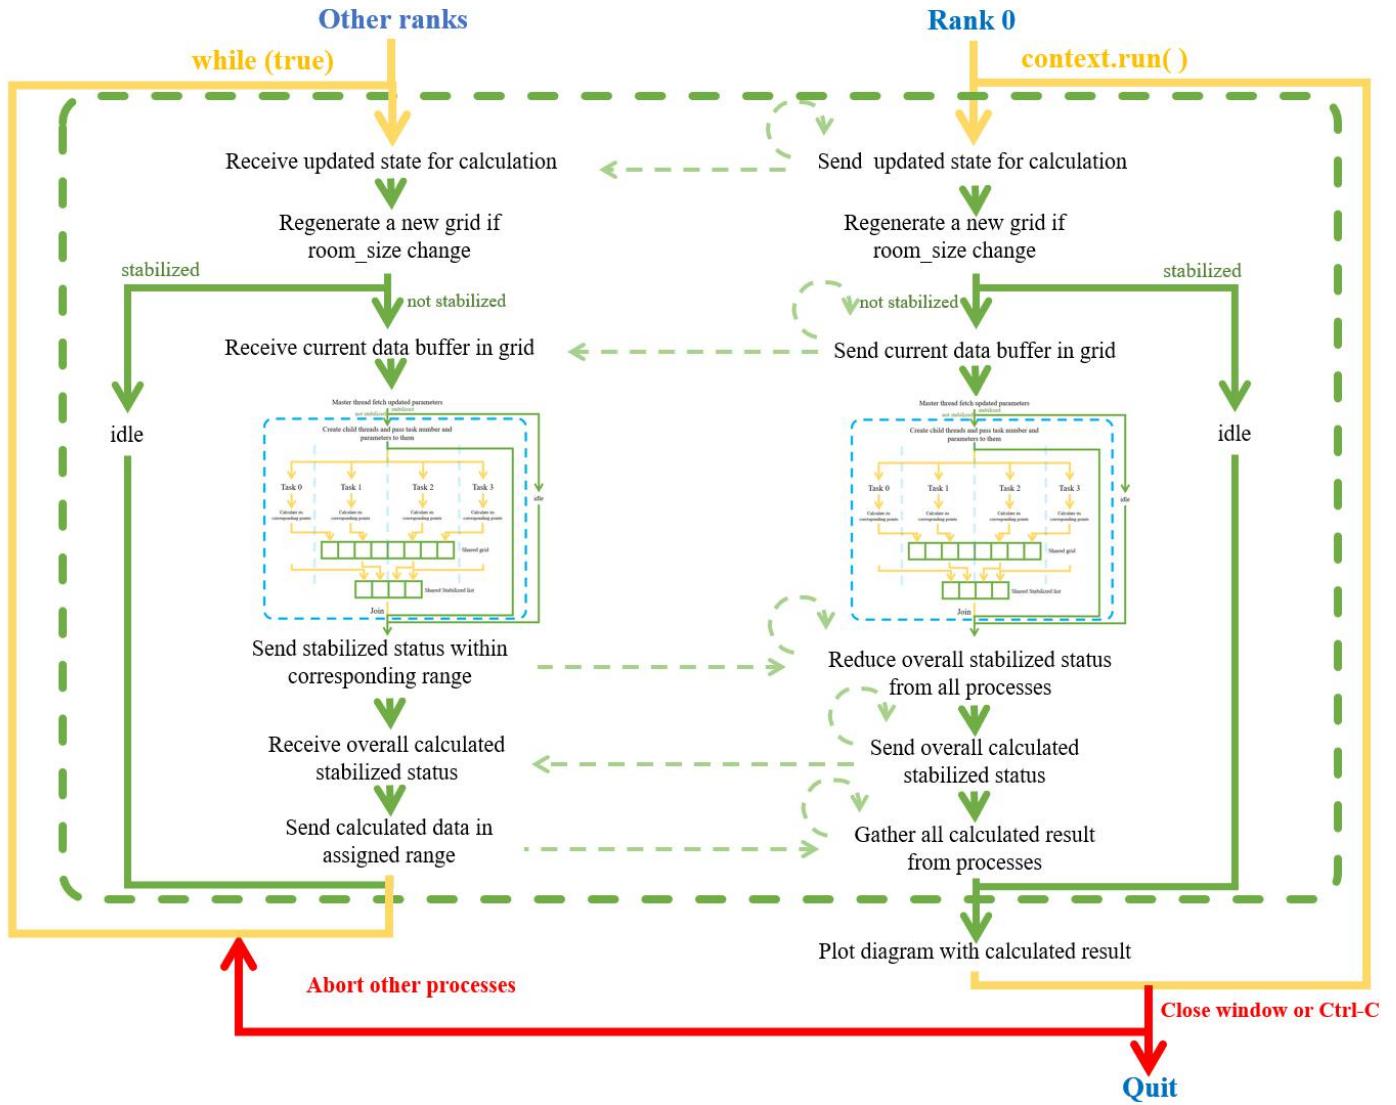
Step 5: All threads do the calculation next temperature on their corresponding regions. After that, they will write their local stabilized status into their corresponding positions on the status list in each process.

Step 6: After joining all the threads, the master thread would then go through all the local status in the status list to generate the overall stabilized status of one process.

Step 4: Rank 0 would reduce the stabilized information of each process local stabilized situation and generate the overall stabilized result which would then broadcast to all other processes.

Step 5: Rank 0 would gather the calculated result from all other processes after which the Rank 0 would plot the heat grid. Next, all process would go to step 1.

Step 6: If GUI window is closed or Ctrl-C is typed in the terminal, the rank 0 process would directly jump to step 7 and abort all other processes. In this case, you may see the abort signal prompted in the terminal which is a normal operation.



3. Program execution

a) Compile and build

First, you need to Put program folder into “/pvfsmnt/(student_id)”. And you need to type “rm -rf build” to remove the previous build folder.

Then follow the procedures to compile and build different version of programs:

i. MPI, Pthread, OpenMPI, MPI with OpenMP

- 1) Type “mkdir build”
- 2) Type “cd build”
- 3) Type “cmake .. -DCMAKE_BUILD_TYPE=Debug”
- 4) Type “cmake --build . -j4”

ii. CUDA

- 1) Type “mkdir build”
- 2) Type “cd build”
- 3) Type “source scl_source enable devtoolset-10”
- 4) Type “CC=gcc CXX=g++ cmake ..”
- 5) Type “make -j12”

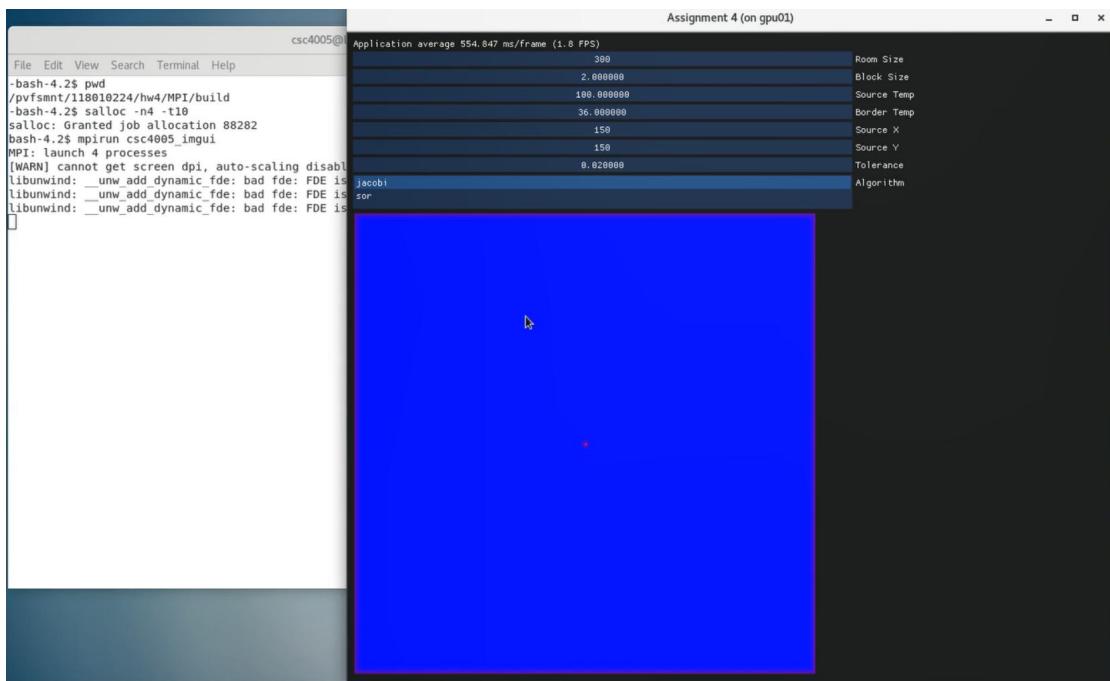
Then, you need to do some authorization settings before executing the programs. In a suitable environment (all the executions of programs in this project are done in the virtual machine set up in

tutorial 1), type “xhost +” first. Then connect to the server using “ssh -Y {student id}@10.26.1.30”. Then type “cd /pvfsmnt/\${whoami}”, “cp ~/.Xauthority /pvfsmnt/\${whoami}”, “export XAUTHORITY=/pvfsmnt/\${whoami}/.Xauthority”.

b) MPI

Since program run with multiple processes, you first need to allocate several cores. Type “salloc -n4 -t10” to request 4 cores for example. Then type “mpirun csc4005_imgui” in the build folder then it will prompt out the program.

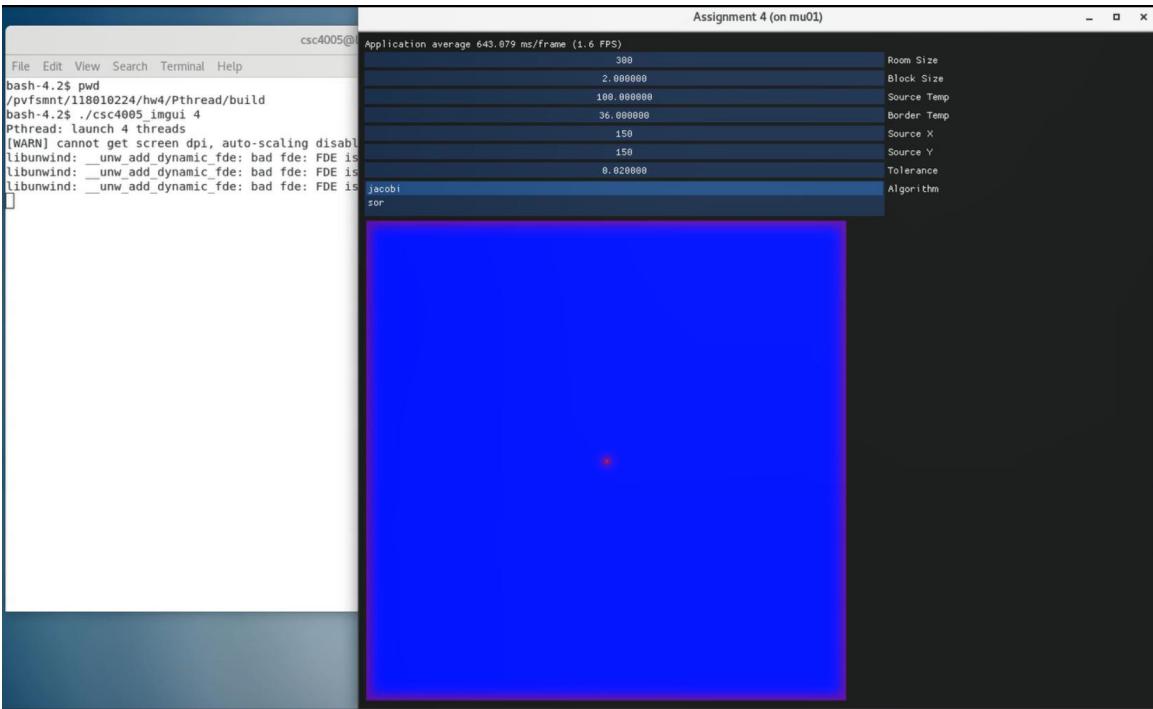
Running result:



c) Pthread

The steps to execute the pthread version program is relative easier, you just need to type “./csc4005_imgui [number of threads]” directly, for example, type “./csc4005_imgui 4” for execute program in 4 threads. The default setting is 4, which would be set up automatically if the input parameters is invalid or not given.

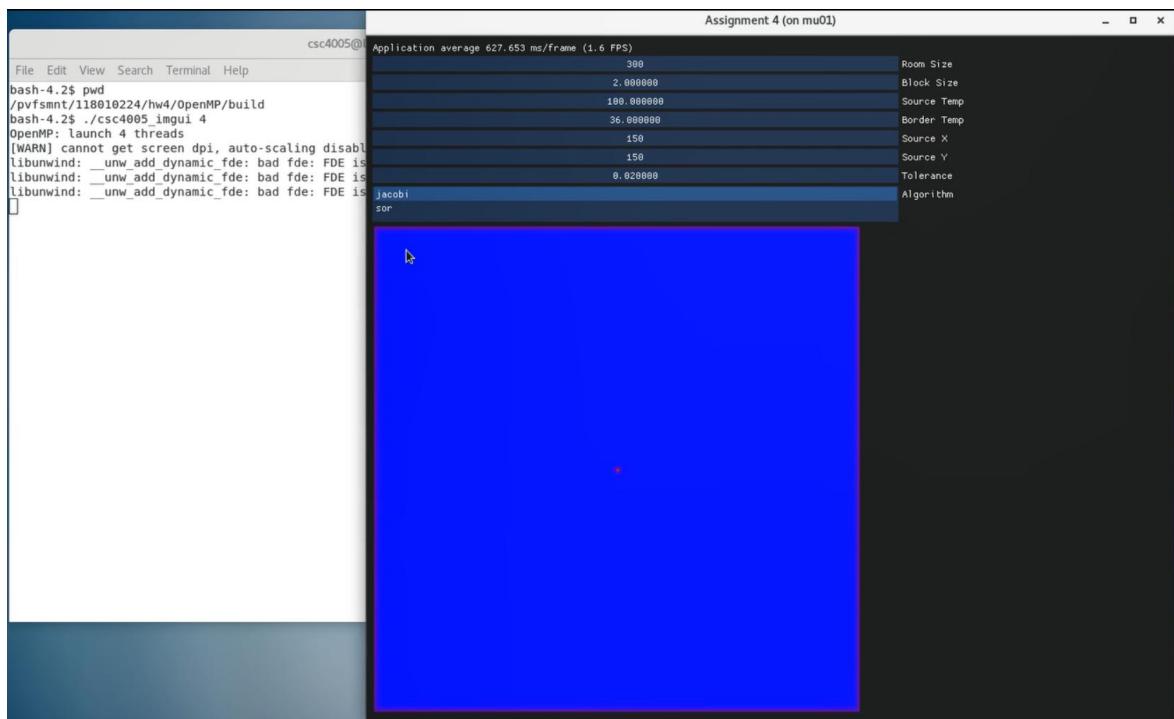
Running result:



d) OpenMP

The steps to execute the OpenMP version program is same as Pthread, you just need to type “./csc4005_imgui [number of threads]” directly, for example, type “./csc4005_imgui 4” for execute program in 4 threads. The default setting is 4, which would be set up automatically if the input parameters is invalid or not given.

Running result:

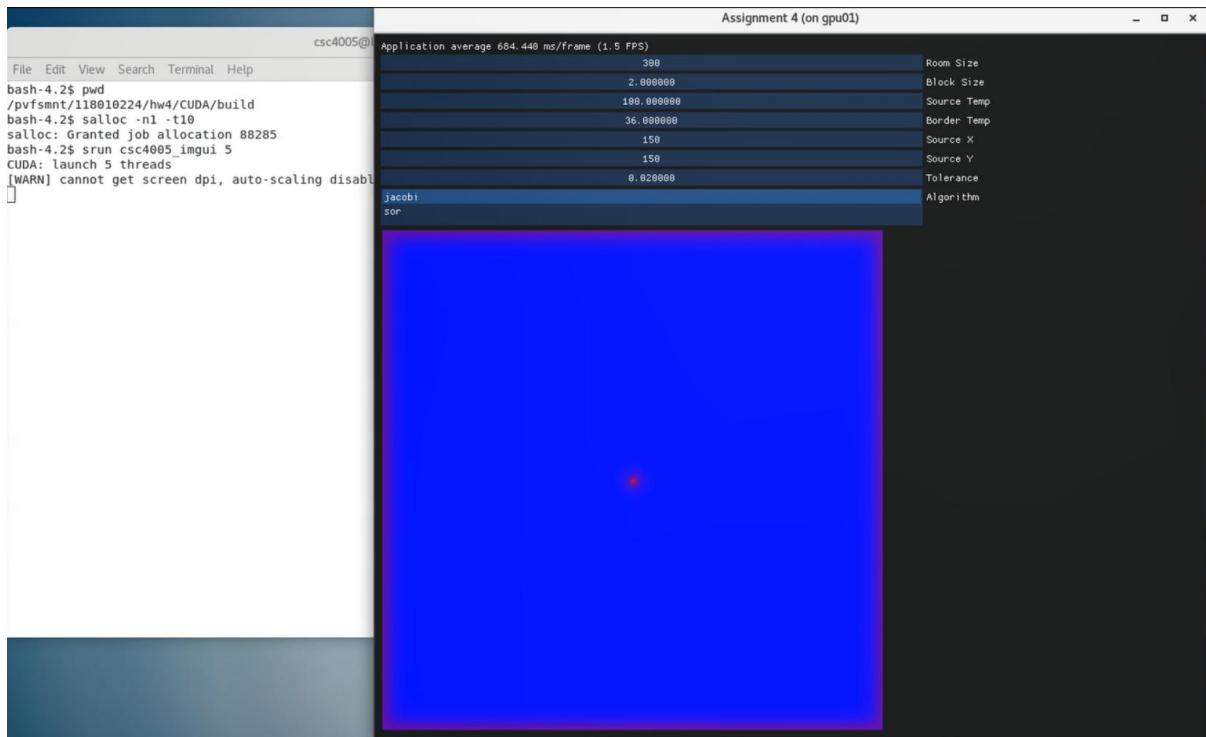


e) CUDA

The steps to execute the OpenMP version program is a little different. You first need to type “`salloc -n1 -t10`” to obtain one GPU node. Then you need to type “`srun ./csc4005_imgui [number of`

threads]” directly, for example, type “./csc4005_imgui 5” for execute program with the CUDA settings as 5 threads in one block. The default setting for all value is 4, which would be set up automatically if the input parameters is invalid or not given.

Running result:

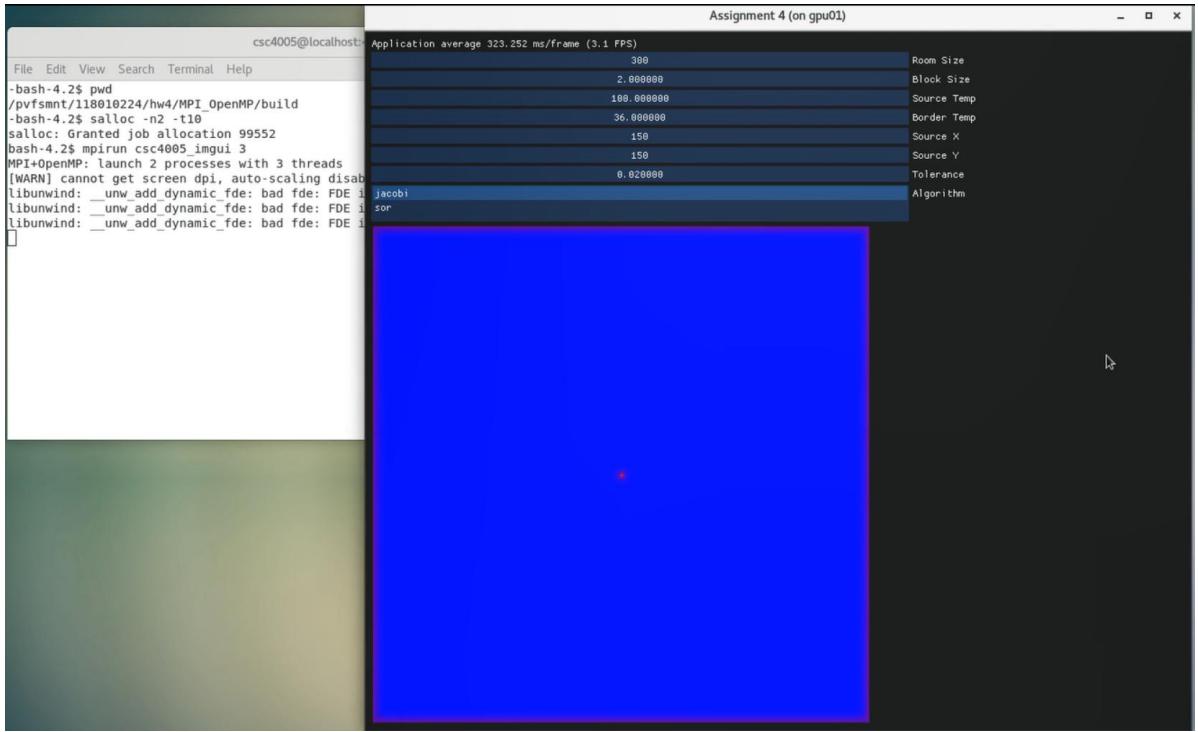


(Notice that the “srun error” in the snapshot is the error of allocating nodes in cluster, which has no relation with the codes)

f) MPI with OpenMP

Since program run with multiple processes, you first need to allocate several cores. Type “`salloc -n2 -t10`” to request 2 cores for example. Then you need to type “`mpirun ./csc4005_imgui [number of threads in each process]`”. The default setting for number of threads in each process is 3, which would be set up automatically if the input parameters is invalid or not given.

Running result:



4. Performance analysis

a) Time complexity

In this analysis, we only consider the calculation task for one iteration. Also, we ignore the plotting steps.

i. Sequential

Assume the total number of points in the grid is N . For Jacobi algorithm, basically, one iteration would simply go through all the points and sum the current values of its 4 neighbors. Overall, it will take $O(4N) = O(N)$ time to finish the job.

ii. MPI

Assume the total number of points in the grid is N and there are p processes. For Jacobi, as for the communication part, the program would take $O(N)$ to broadcast its current data buffer and $O(p)$ to reduce the stabilized results in every process and $O(1)$ to broadcast the reduced result. Finally, it will take $O(N)$ to gather result. Therefore, we have $T_{comm} = O(N) + O(p) + O(1) + O(N) = O(N)$. Then for computation, each process would only need to take care of $\frac{N}{p}$ points which means $T_{comp} = O\left(4 * \frac{N}{p}\right) = O\left(\frac{N}{p}\right)$. Finally, the result would be $T = T_{comm} + T_{comp} = O(N) + O\left(\frac{N}{p}\right) = O(N)$. For Sor, the situation is nearly the same which has same time complexity as $O(N)$.

iii. Pthread/OpenMP/CUDA

Assume the total number of points in the grid is N and there are t threads. For Jacobi, first, it may take $O(t)$ to create t threads. With t threads running in parallel, each thread would only need to take $O\left(\frac{N}{t}\right)$ to do the calculation. Finally, master thread would take $O(t)$ to check all

stabilized situation in each thread to generate the overall stabilized status. Overall, we have $T = O\left(\frac{N}{t}\right) + O(t)$. The time complexity for Sor is quite similar which is also $T = O\left(\frac{N}{t}\right) + O(t)$. For CUDA, the time complexity for Jacobi and Sor would be $T = O\left(\frac{N}{t}\right) + O(t) + O(N) = O(N)$ since it would take extra time to copy and update the data buffer in grid.

iv. MPI with OpenMP

Assume the total number of points in the grid is N and there are p processes and each process create t threads. In this case, we just replace the computation time in MPI with the time complexity of OpenMP then we get the overall time complexity which is $T = T_{comm} + T_{comp} = O(N) + O\left(\frac{N}{pt}\right) + O(t) + O(N) = O(N)$.

b) Data race test in Pthread and OpenMP

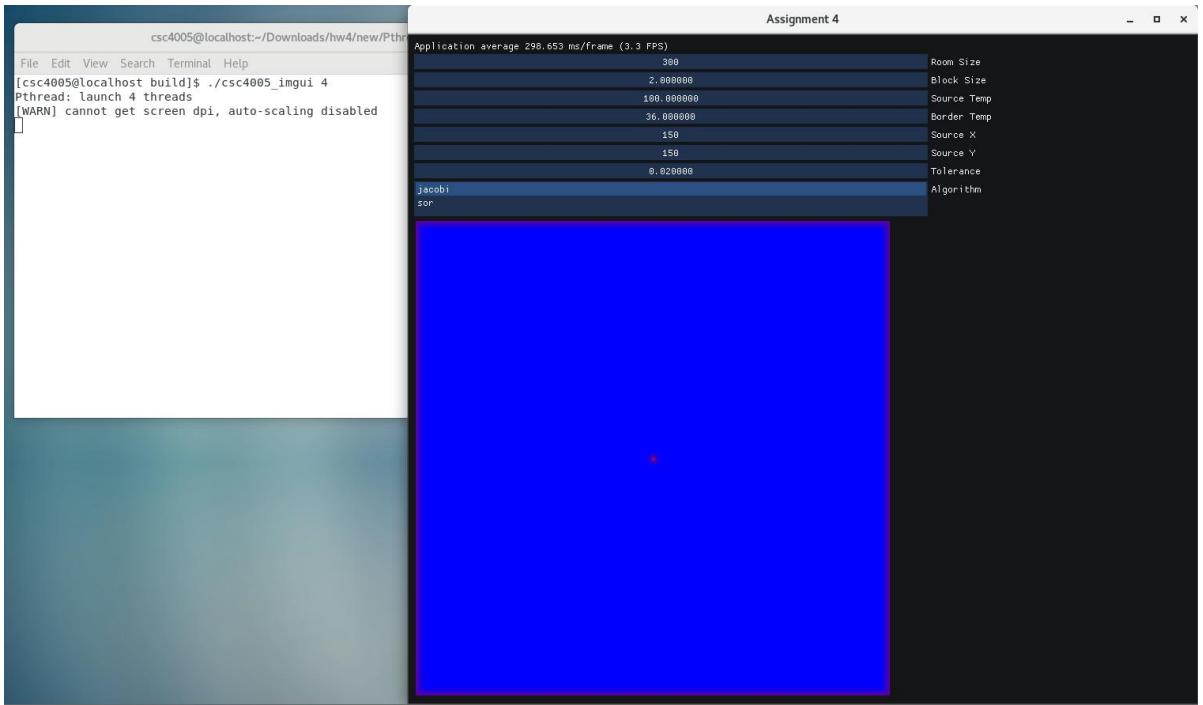
Notice that all data race tests are conducted on virtual machine provided in the tutorials.

i. Pthread

To activate the data race mode, you need to do some modifications in the CMakeList.txt, which is shown in the followings:

```
Pthread > CMakeLists.txt
1 cmake_minimum_required(VERSION 3.2)
2 project(csc4005_imgui)
3
4 find_package(SDL2 REQUIRED)
5 find_package(Freetype REQUIRED)
6 find_package(MPI REQUIRED)
7 find_package(Threads REQUIRED)
8 set(CMAKE_CXX_STANDARD 20)
9 set(CMAKE_CXX_EXTENSIONS ON)
10 set(OpenGL_GL_PREFERENCE "GLVND")
11
12 set(CMAKE_CXX_FLAGS "-fsanitize=thread")
13
14 find_package(OpenGL REQUIRED)
15
16 include_directories(
17     include
18     imgui
19     imgui/backends
20     ${SDL2_INCLUDE_DIRS}
21     ${FREETYPE_INCLUDE_DIRS}
22     ${MPI_CXX_INCLUDE_DIRS})
23
24 file(GLOB IMGUI_SRC
25     imgui/*.cpp
26     imgui/backends/imgui_impl_sdl.cpp
27     imgui/backends/imgui_impl_opengl2.cpp
28     imgui/misc/freetype/imgui_freetype.cpp
29     imgui/misc/cpp/imgui_stl.cpp
30 )
31 add_library(core STATIC ${IMGUI_SRC})
32 file(GLOB CSC4005_PROJECT_SRC src/*.cpp src/*.c)
33 add_executable(csc4005_imgui ${CSC4005_PROJECT_SRC})
34 get_filename_component(FONT_PATH imgui/misc/fonts/DroidSans.ttf ABSOLUTE)
35 target_link_libraries(core PUBLIC
36     Freetype::Freetype SDL2::SDL2 OpenGL::GL ${CMAKE_DL_LIBS} Threads::Threads ${MPI_CXX_LIBRARIES})
```

Running result:



ii. OpenMP

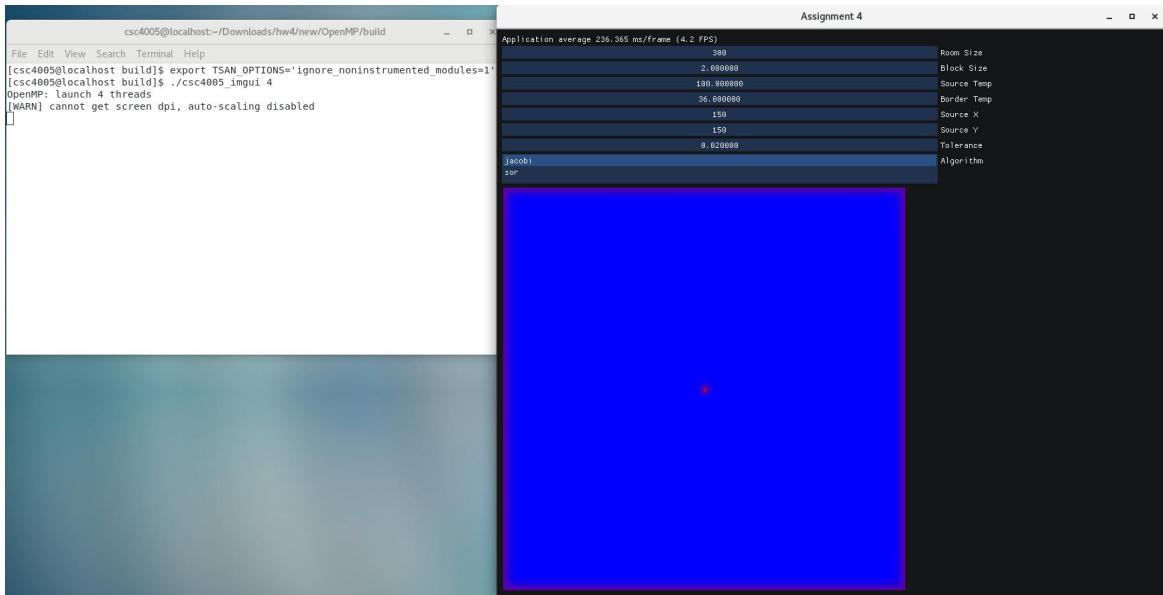
First, you need to do some modifications in CMakeList.txt before compile and build, which is shown in the followings:

```
CMakeLists.txt
OpenMP > CMakeLists.txt
4   find_package(SDL2 REQUIRED)
5   find_package(Freetype REQUIRED)
6   find_package(MPI REQUIRED)
7   find_package(Threads REQUIRED)
8   set(CMAKE_CXX_STANDARD 20)
9   set(CMAKE_CXX_EXTENSIONS ON)
10  set(OpenGL_GL_PREFERENCE "GLVND")
11  find_package(OpenGL REQUIRED)
12
13  find_package(OpenMP REQUIRED)
14  set (CMAKE_C_FLAGS "${CMAKE_C_FLAGS} ${OpenMP_C_FLAGS}")
15  # set (CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS} ${OpenMP_CXX_FLAGS}")
16  set (CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS} ${OpenMP_CXX_FLAGS} -fsanitize=thread")
17  set (CMAKE_EXE_LINKER_FLAGS "${CMAKE_EXE_LINKER_FLAGS} ${OpenMP_EXE_LINKER_FLAGS}")
18
19
20  include_directories(
21    include
22    imgui
23    imgui/backends
24    ${SDL2_INCLUDE_DIRS}
25    ${FREETYPE_INCLUDE_DIRS}
26    ${MPI_CXX_INCLUDE_DIRS})
27
28  file(GLOB IMGUI_SRC
29    imgui/*.cpp
30    imgui/backends/imgui_impl_sdl.cpp
31    imgui/backends/imgui_impl_opengl2.cpp
32    imgui/misc/freetype/imgui_freetype.cpp
33    imgui/misc/cpp/imgui_stl.cpp
34    )
35  add_library(core STATIC ${IMGUI_SRC})
36  file(GLOB CSC4005_PROJECT_SRC src/*.cpp src/*.c)
37  add_executable(csc4005 imgui ${CSC4005_PROJECT_SRC})
38  get_filename_component(FONT_PATH imgui/misc/fonts/DroidSans.ttf ABSOLUTE)
39  target_link_libraries(core PUBLIC
40    Freetype::Freetype SDL2::SDL2 OpenGL::GL ${CMAKE_DL_LIBS} Threads::Threads ${MPI_CXX_LIBRARIES})
```

Before running the program, you need to type “ export

TSAN_OPTIONS='ignore_noninstrumented_modules=1' ” to avoid some extra false positive errors.

Running result:



c) Performance Test

In this project, MPI, Pthread, OpenMP, CUDA and MPI with OpenMP would work in the same way as sequential version if total number of process or thread is 1. Thus, for testing, all versions with one process would work as the sequential version in their test group. All the test results generated from all versions would only focus on the calculation steps. Notice that we only consider one iteration of calculation instead of waiting the grid to be stabilized since all calculation are almost same in each iteration.

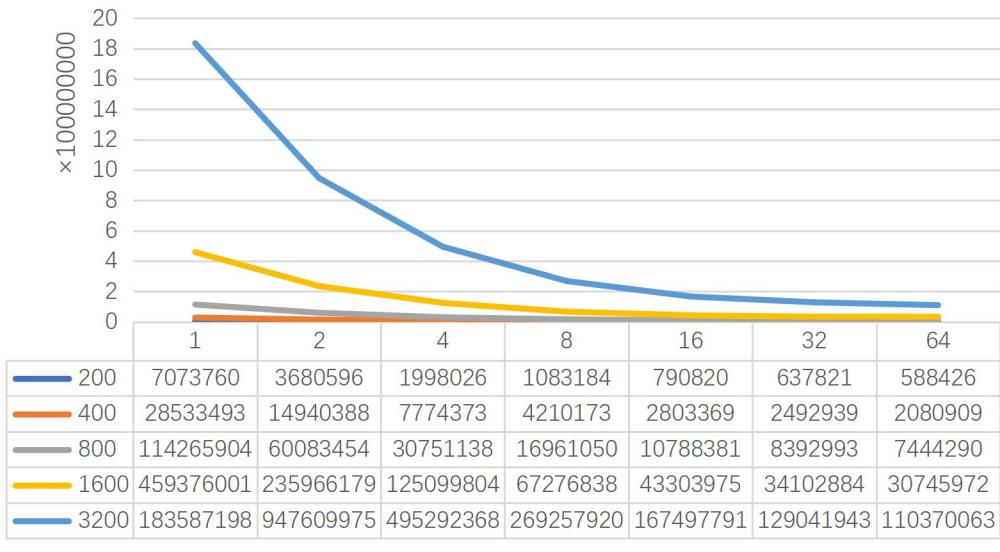
i. Time vs number of Core/Thread

Notice for testing, MPI will use one node for 1 process test and 2 nodes for all the remaining tests. For CUDA, we will only use one node. For MPI with OpenMP, we will only use one node for totally one thread test and two nodes for the remaining test.

Part 1. All test sets:

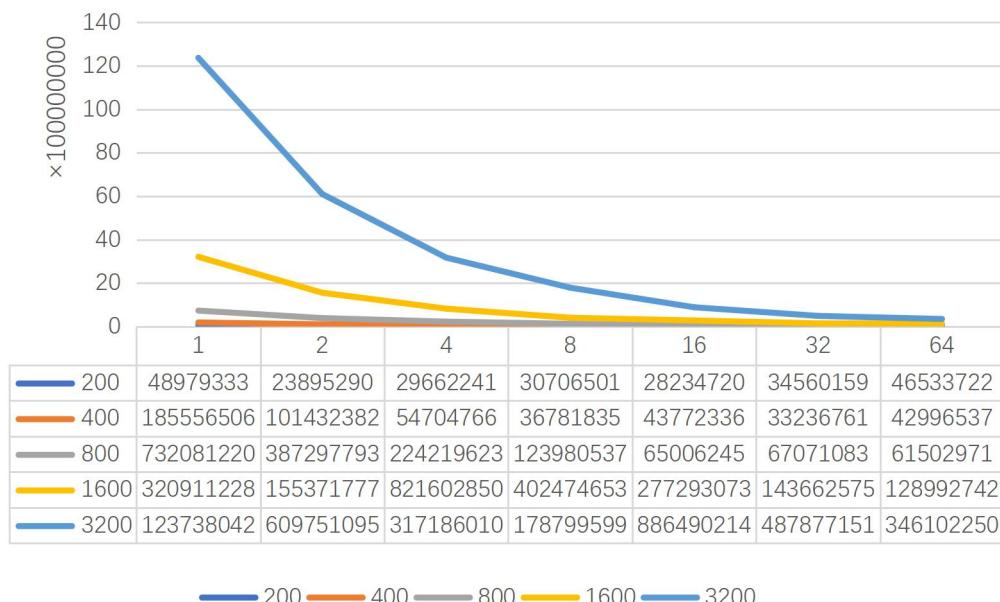
MPI

Time(ms) vs Core Number



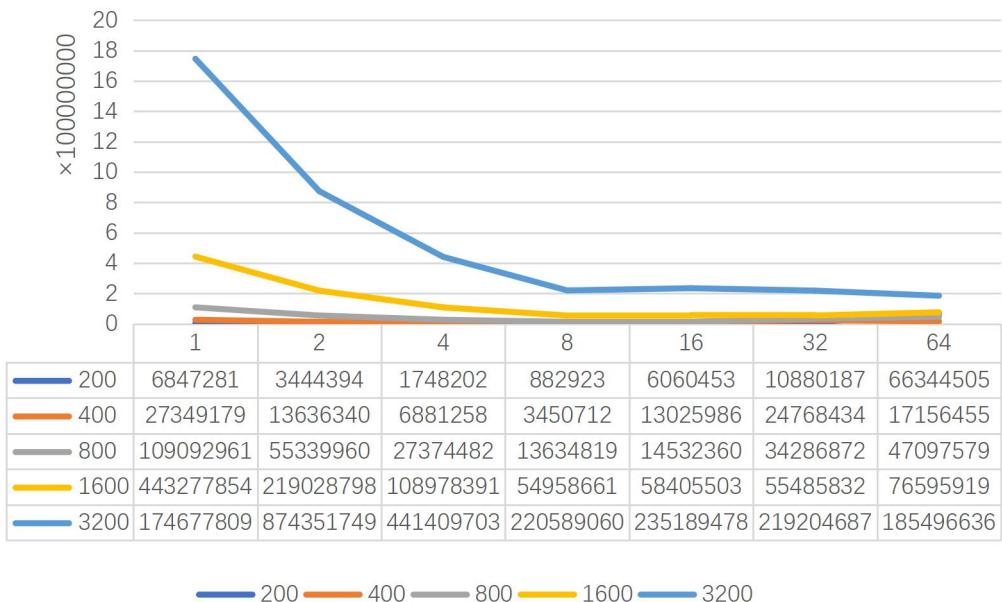
Pthread

Time(ms) vs Core Number



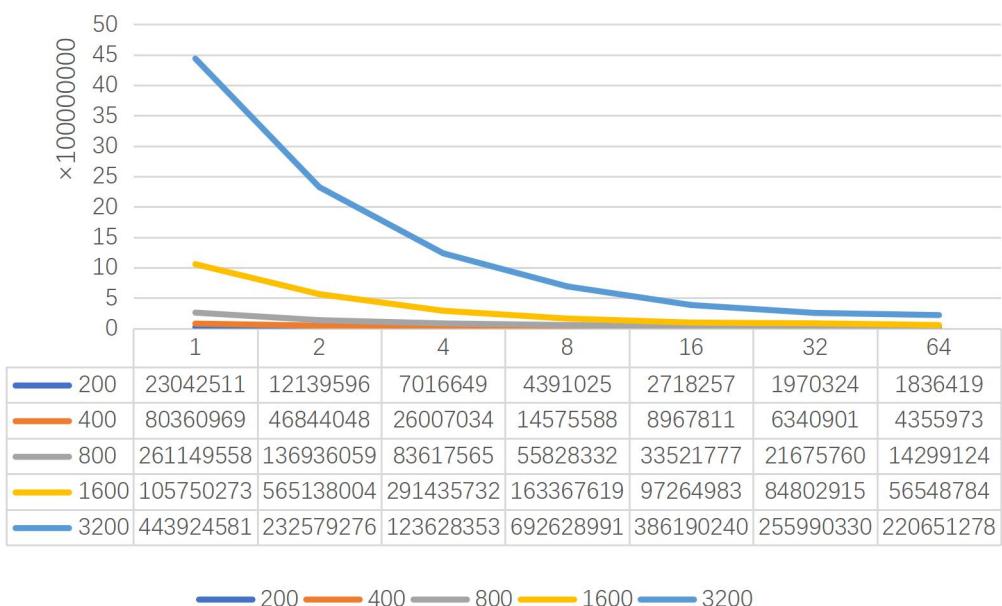
OpenMP

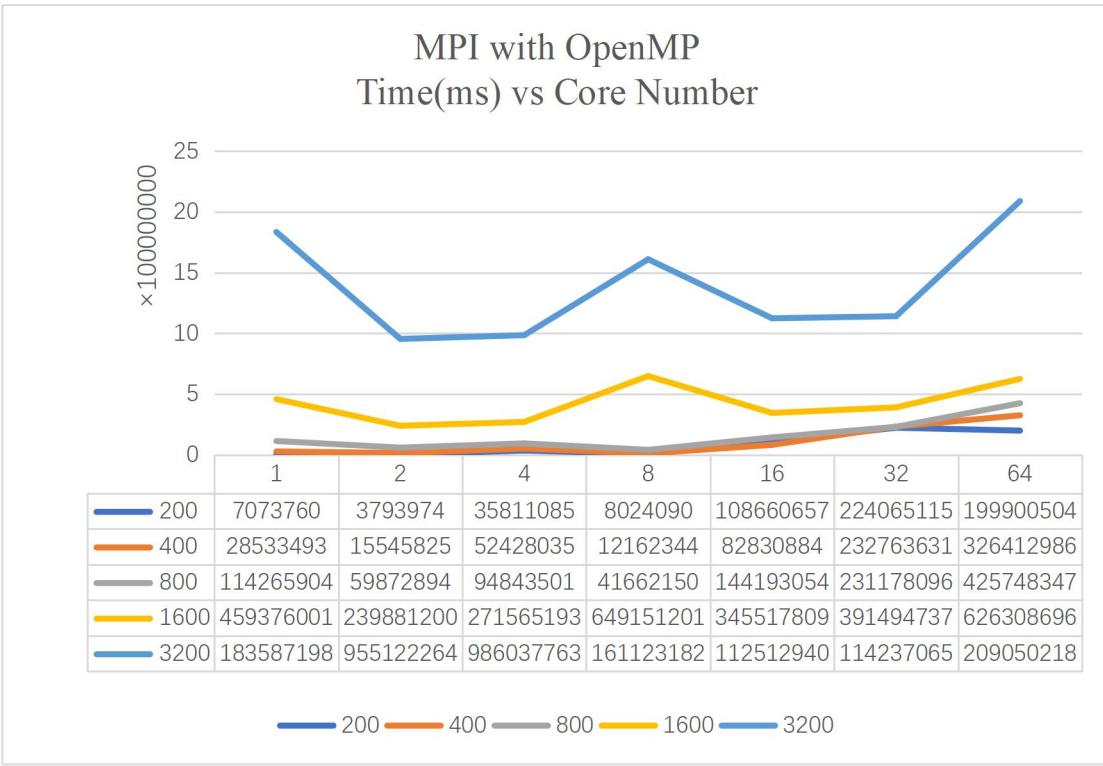
Time(ms) vs Core Number



CUDA

Time(ms) vs Core Number

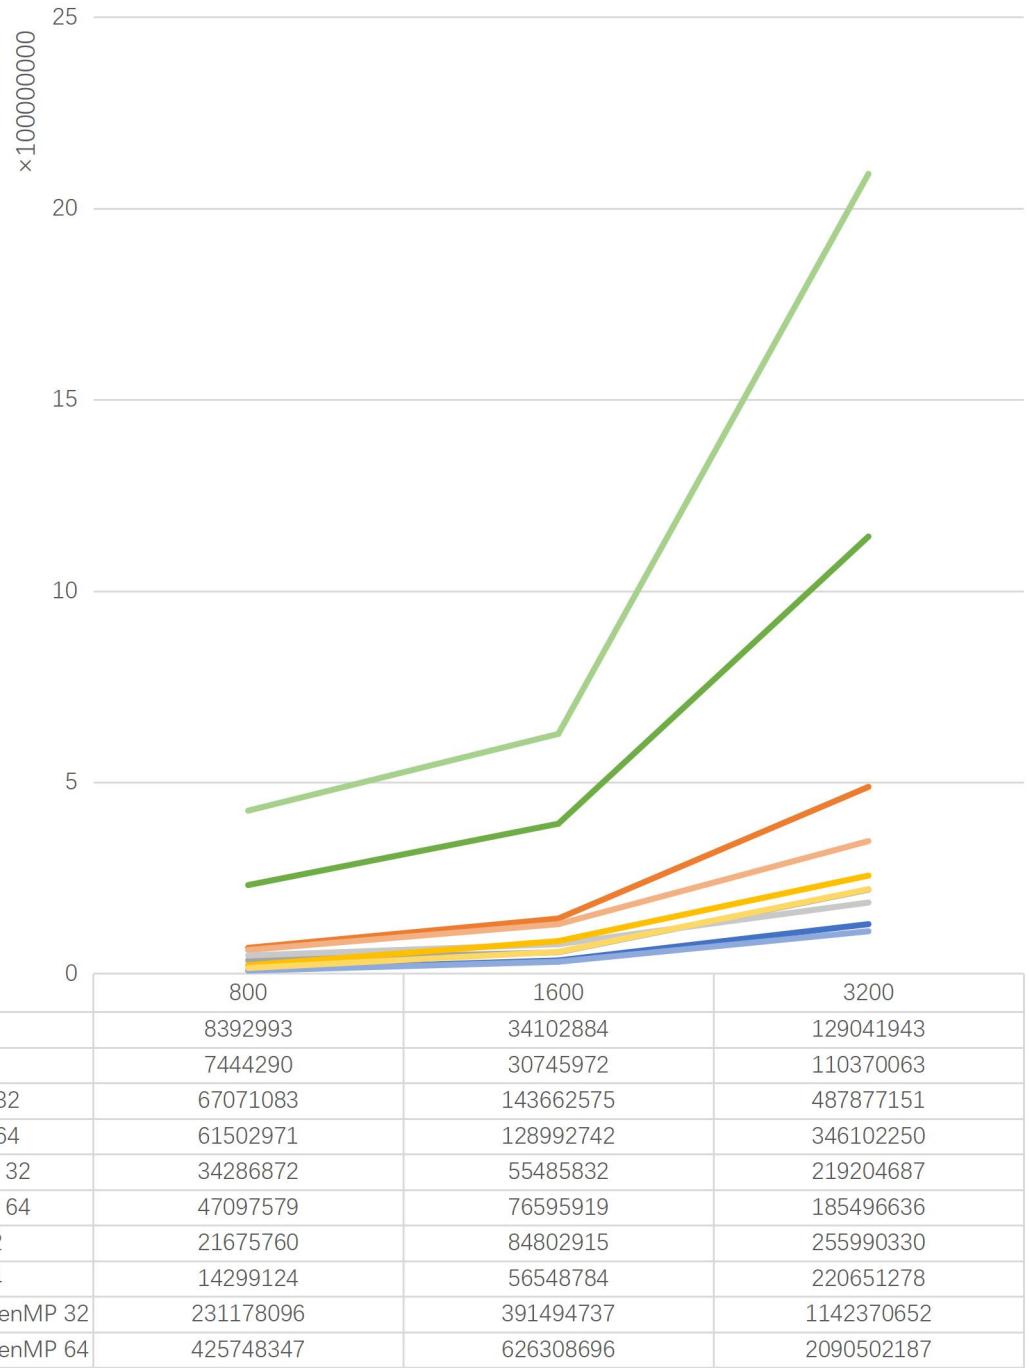




Generally, with more processes or threads involves, the execution time are further decreasing. Except MPI with OpenMP, the testing result shows that double the number of processes or threads would decrease the execution time by half, which indicating the programs are cost optimum. Things are a little different in MPI with OpenMP where the running time fluctuate with the increase of threads. One of the possible reasons is the communication overhead and threads creation effect offset the speedup bring up by parallel computing. Also, we can see that a good performance was generally reached when we have relatively small number of threads.

Part 2. All versions in larger body test size and large process/thread size

All versions in larger body test size and large process/thread size



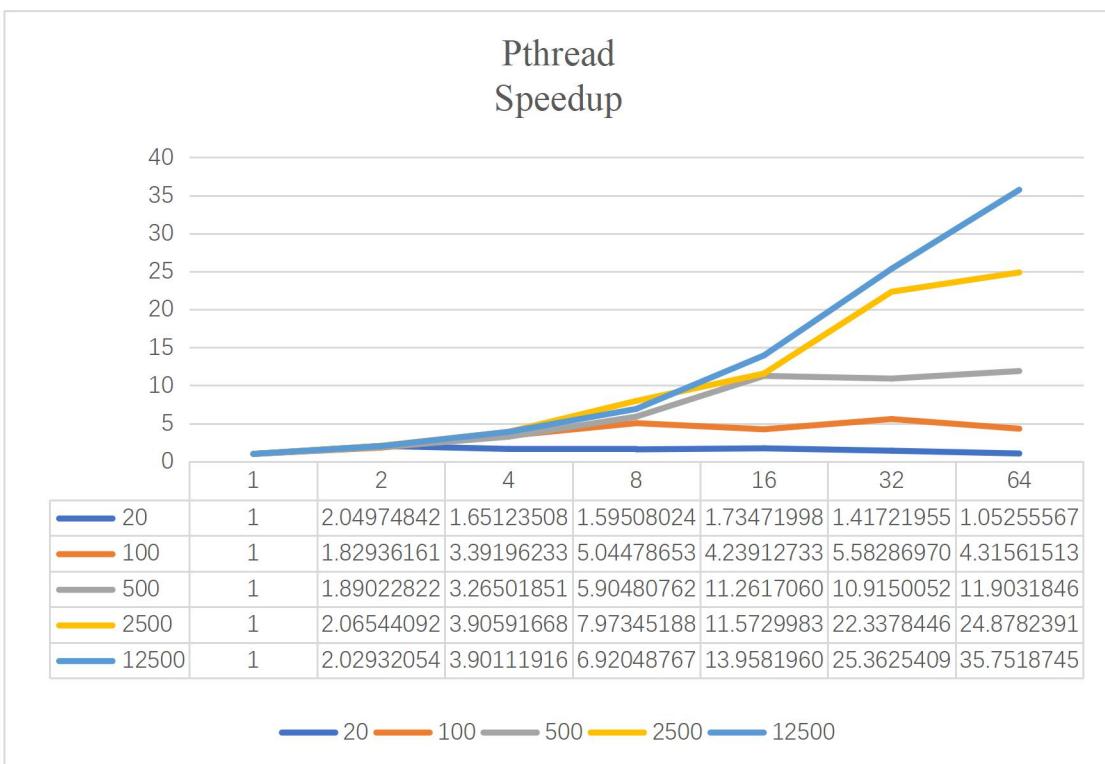
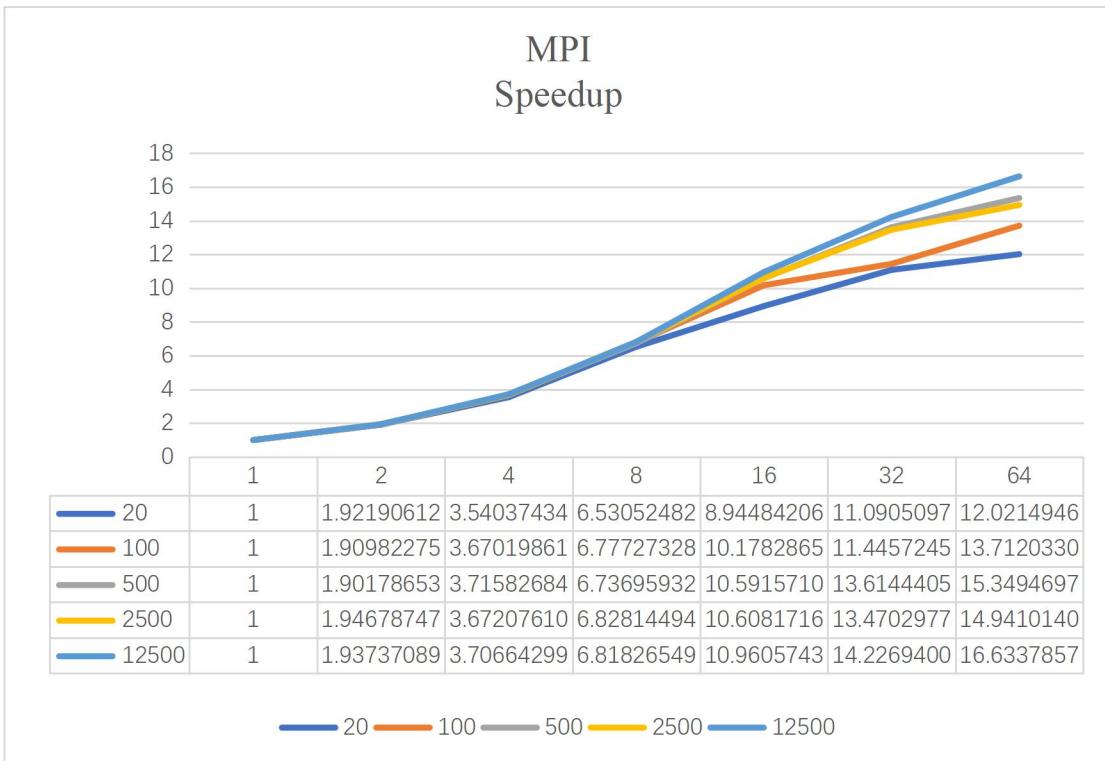
In the test with large data size and large thread or process size, the MPI with OpenMP generally works the worst among all which may be attributed to the communication time and the threads creation and management time. Among all, MPI performs the best, one of the possible reasons may be the relative powerful computing ability of each process.

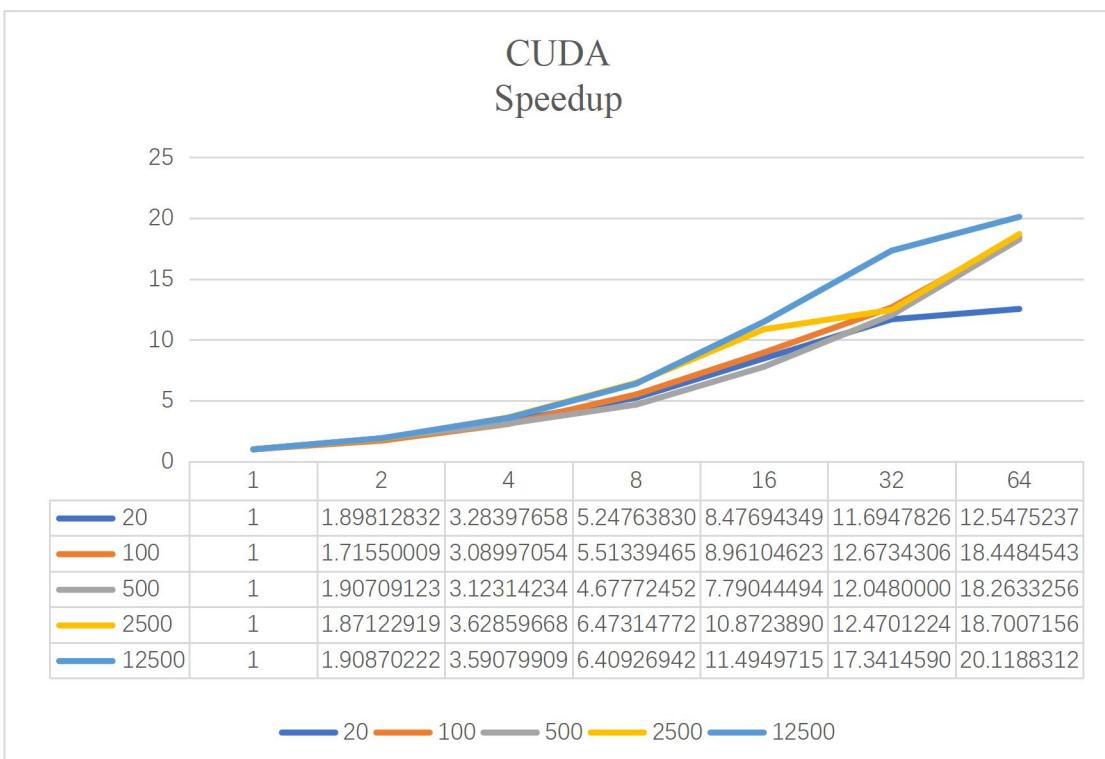
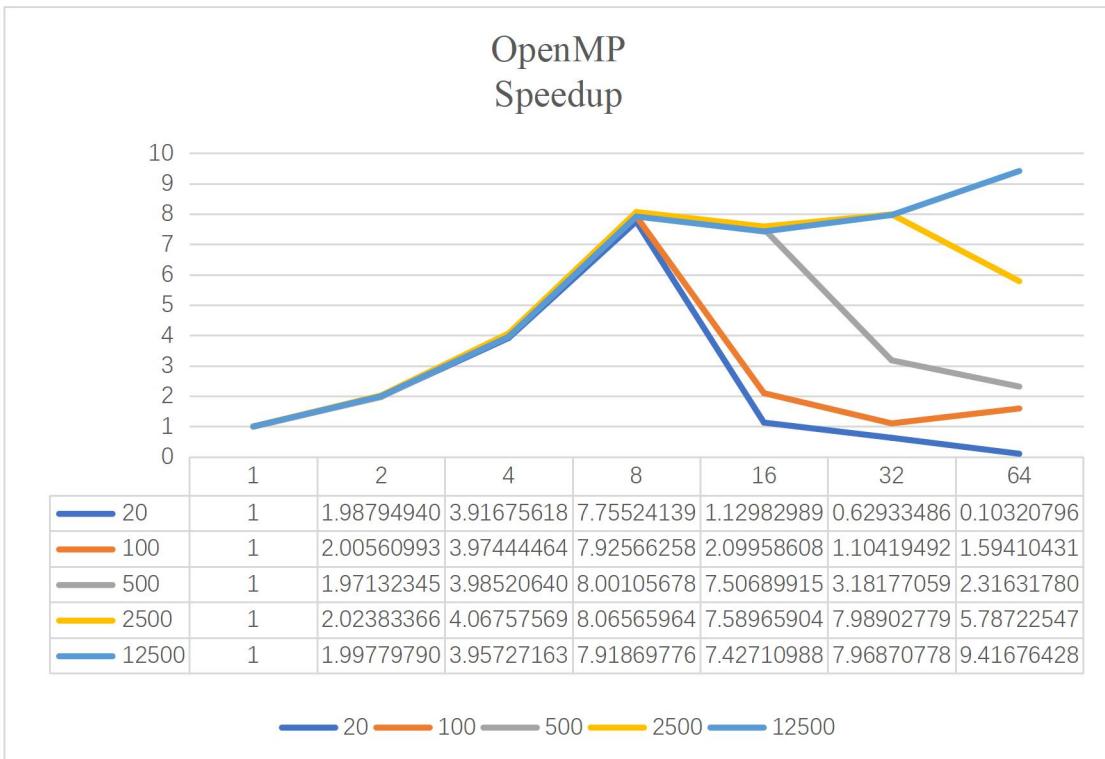
Part 3. Speedup

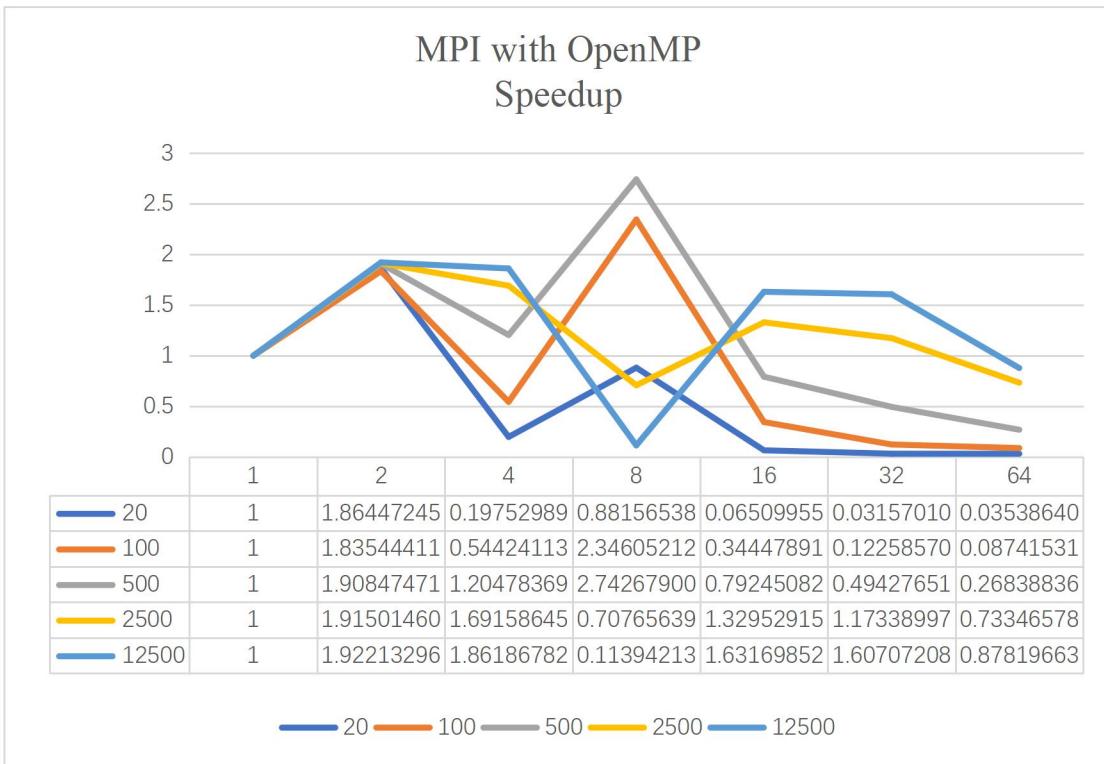
We can also calculate the Speedup of the program.

The formula of the Speedup is given as:

$$\text{Speedup} = \frac{\text{running time of sequential version}}{\text{running time of parallel version}}$$





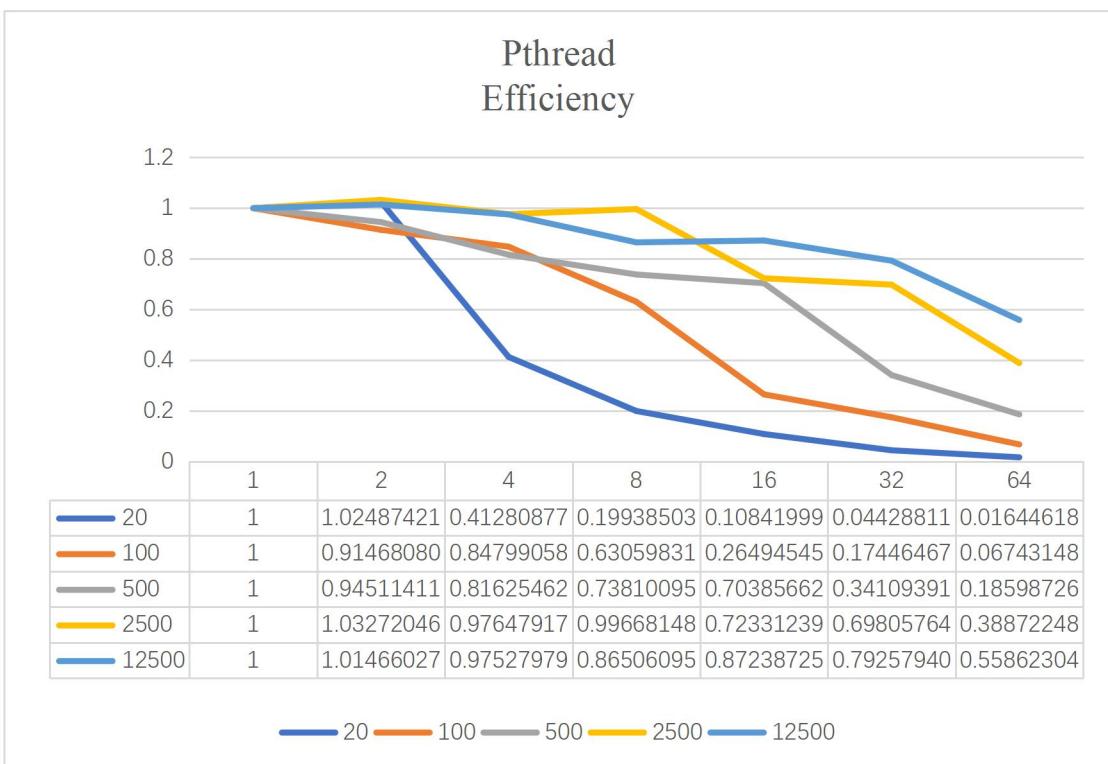
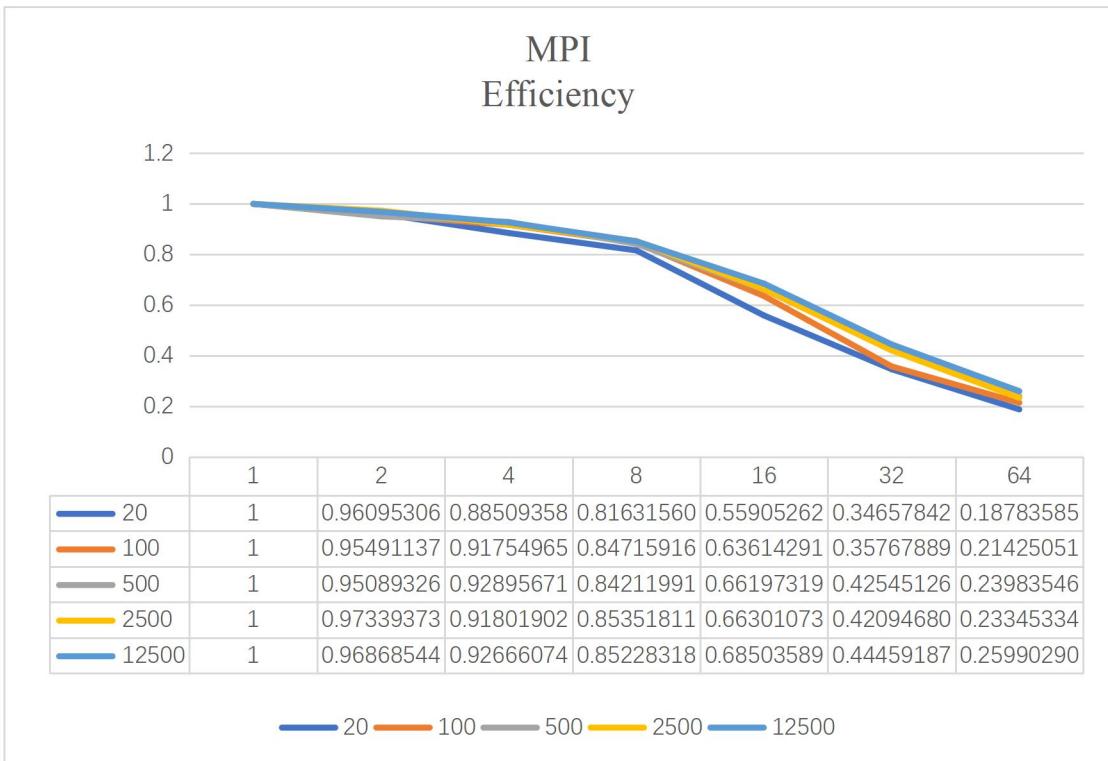


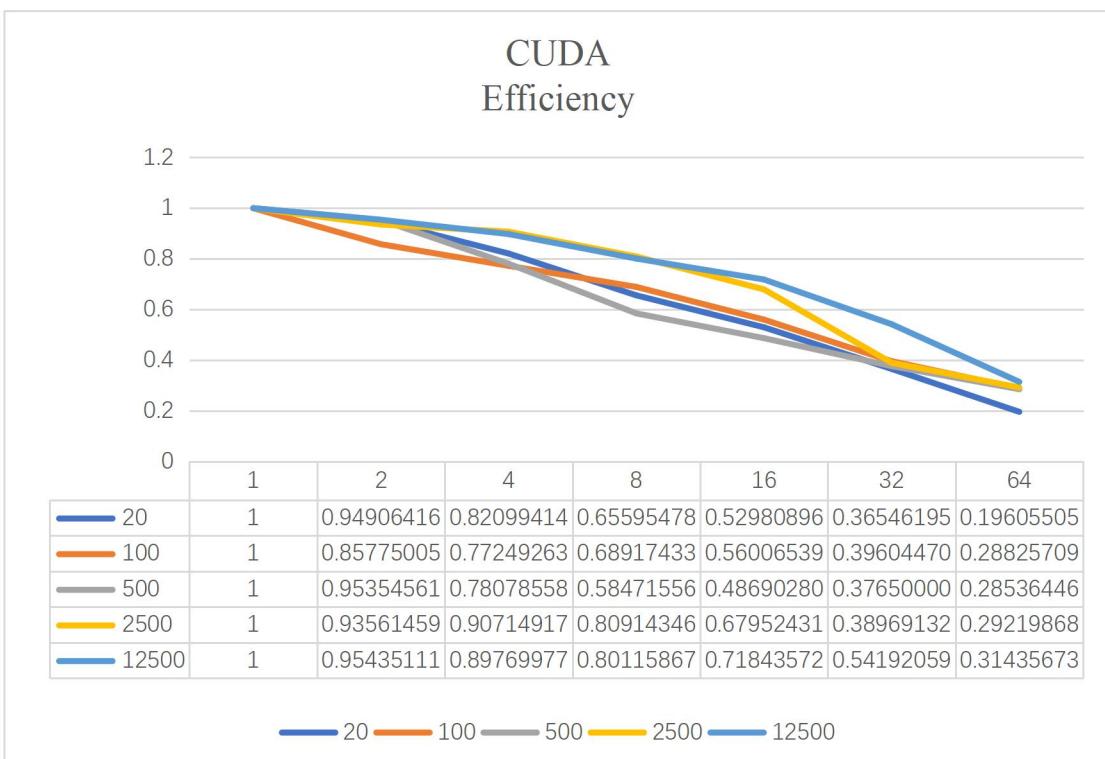
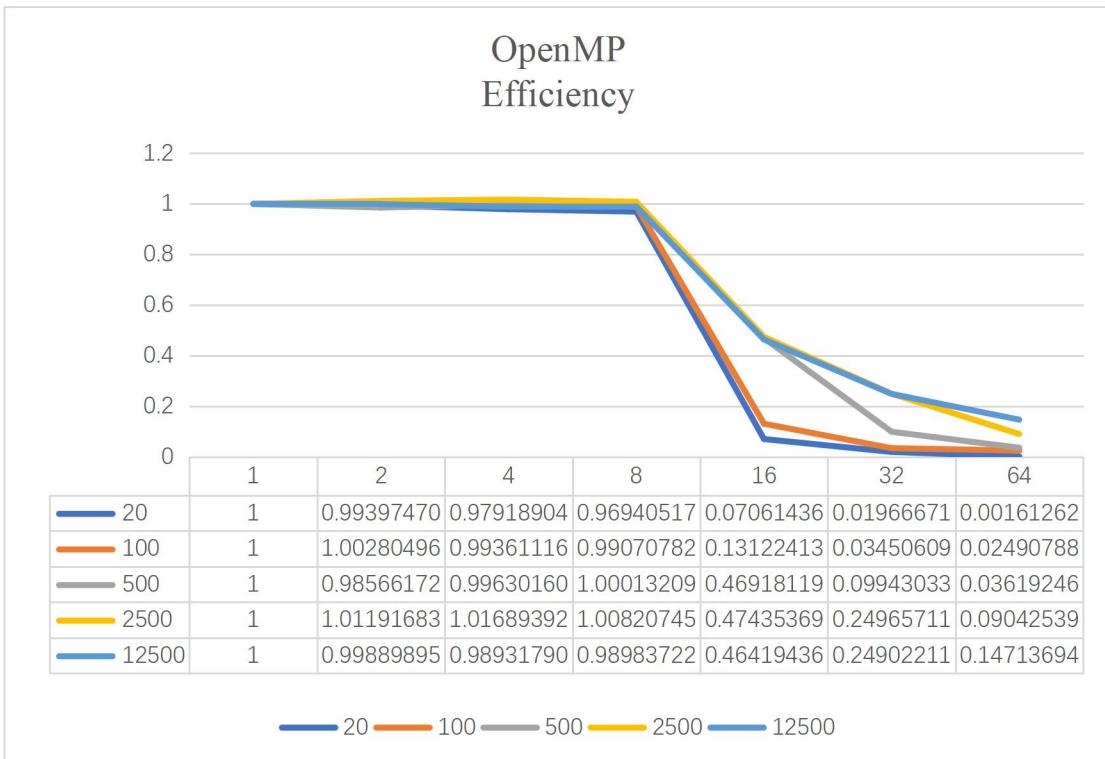
Generally, all settings except the MPI with OpenMP achieve the best speedup in the large data size and large number of processes or threads. In this case, the communication overhead and threads management overhead are far much smaller than the time spent on computation. For the OpenMP, we can easily see that the performance are the worst in the smallest data test size which is definitely resulted from the extra threads management time. The speedup of MPI with OpenMP are quite different, one possible reason is the communication overhead and threads creation effect offset the speedup bring up by parallel computing.

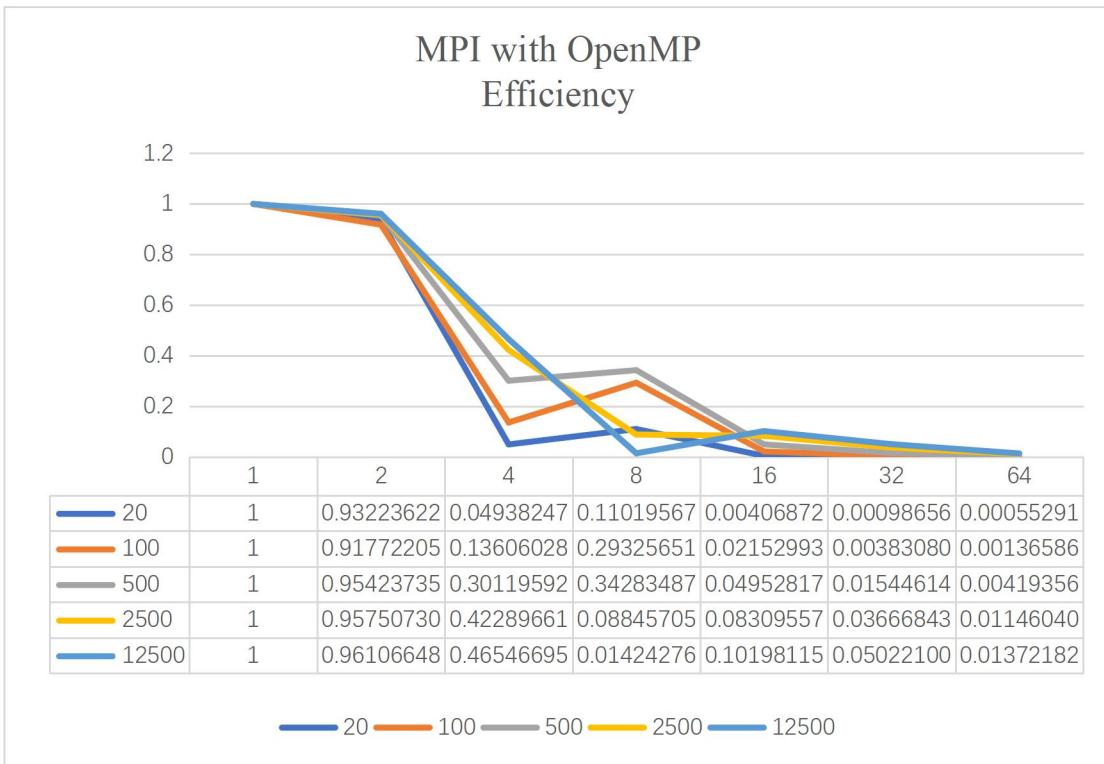
Part 4. Efficiency

The efficiency formula is given as:

$$E = \frac{\text{Speedup}}{\text{number of processes(or threads)}}$$







Generally, the efficiency would decrease with the increase of size of processes and threads since with more threads or process, each process would be assigned with less subtask size which may spent relatively small time for computing and more time for waiting other processes or threads to finish their tasks. For Pthread, when handling large data size with small number of threads, it could sometimes perform quite good at efficiency which is slightly larger than 1. The similar situation was also witness in OpenMP, when using smaller number of threads, it could still get a good efficiency which is remained around 1.

5. Conclusion and Limitations

This report mainly discusses five parallel versions (MPI, Pthread, OpenMP, CUDA, MPI with OpenMP) of the given sequential heat distribution sequential program. For parallel version, the design of program considers the potential data race and avoid it. Tests with different number of cores and threads and data size, the report shows the improvement of parallel calculation and also discuss the communication overhead and thread management overhead may offset the performance improvement brought by parallel calculation especially in MPI with OpenMP.

To further improve the computation efficiency in MPI, we can actually use scatter to distribute the data buffer of grid instead of using broadcast directly. Notice that when do the calculation of every point, we only need 4 neighbor points' information. In this scenario we could decrease the communication overhead of data distributing from $O(N)$ to $O\left(\frac{N}{p}\right)$ which could reduce the useless data transformation time and computing resources.

Codes:

1) MPI

a) main.cpp

```
#include <graphic/graphic.hpp>
#include <imgui_impl_sdl.h>
#include <cstring>
#include <chrono>
#include <hdist/hdist.hpp>

template<typename ...Args>
void UNUSED(Args &&... args [[maybe_unused]]) {}

ImColor temp_to_color(double temp) {
    auto value = static_cast<uint8_t>(temp / 100.0 * 255.0);
    return {value, 0, 255 - value};
}

static MPI_Datatype MPI_State;

int main(int argc, char **argv) {
    int rank;
    int processSize;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &processSize);

    // initialize MPI struct
    const int nitems=9;
    int blocklengths[9] = {1,1,1,1,1,1,1,1,1};
    MPI_Datatype types[9] = {MPI_INT, MPI_FLOAT, MPI_INT, MPI_INT,
                           MPI_FLOAT, MPI_FLOAT, MPI_FLOAT, MPI_FLOAT, MPI_INT};

    MPI_Aint offsets[9];

    offsets[0] = offsetof(hdist::State, room_size);
    offsets[1] = offsetof(hdist::State, block_size);
    offsets[2] = offsetof(hdist::State, source_x);
    offsets[3] = offsetof(hdist::State, source_y);
    offsets[4] = offsetof(hdist::State, source_temp);
    offsets[5] = offsetof(hdist::State, border_temp);
    offsets[6] = offsetof(hdist::State, tolerance);
    offsets[7] = offsetof(hdist::State, sor_constant);
    offsets[8] = offsetof(hdist::State, algo);

    MPI_Type_create_struct(nitems, blocklengths, offsets, types, &MPI_State);
    MPI_Type_commit(&MPI_State);
```

```

UNUSED(argc, argv);
bool first = true;
bool finished = false;
static hdist::State current_state, last_state;
auto grid = hdist::Grid{
    static_cast<size_t>(current_state.room_size),
    current_state.border_temp,
    current_state.source_temp,
    static_cast<size_t>(current_state.source_x),
    static_cast<size_t>(current_state.source_y)};

if(rank == 0){
    printf("MPI: launch %d processes\n", processSize);

    // test codes
    // int testCases[] = {200, 400, 800, 1600, 3200};
    // int testCaseIdx = 0, testCnt = 0;
    // static std::chrono::high_resolution_clock::time_point beginIteration,
endIteration;

    static std::chrono::high_resolution_clock::time_point begin, end;
    static const char* algo_list[2] = { "jacobi", "sor" };
graphic::GraphicContext context{"Assignment 4"};
context.run([&](graphic::GraphicContext *context [[maybe_unused]], SDL_Window *) {
    auto io = ImGui::GetIO();
    ImGui::SetNextWindowPos(ImVec2(0.0f, 0.0f));
    ImGui::SetNextWindowSize(io.DisplaySize);
    ImGui::Begin("Assignment 4", nullptr,
        ImGuiWindowFlags_NoMove
        | ImGuiWindowFlags_NoCollapse
        | ImGuiWindowFlags_NoTitleBar
        | ImGuiWindowFlags_NoResize);
    ImDrawList *draw_list = ImGui::GetWindowDrawList();
    ImGui::Text("Application average %.3f ms/frame (%.1f FPS)", 1000.0f /
    ImGui::GetIO().Framerate,
        ImGui::GetIO().Framerate);
    ImGui::DragInt("Room Size", &current_state.room_size, 10, 200, 1600, "%d");
    ImGui::DragFloat("Block Size", &current_state.block_size, 0.01, 0.1, 10, "%f");
    ImGui::DragFloat("Source Temp", &current_state.source_temp, 0.1, 0, 100, "%f");
    ImGui::DragFloat("Border Temp", &current_state.border_temp, 0.1, 0, 100, "%f");
    ImGui::DragInt("Source X", &current_state.source_x, 1, 1, current_state.room_size
- 2, "%d");
    ImGui::DragInt("Source Y", &current_state.source_y, 1, 1, current_state.room_size
- 2, "%d");
    ImGui::DragFloat("Tolerance", &current_state.tolerance, 0.01, 0.01, 1, "%f");
    ImGui::ListBox("Algorithm", reinterpret_cast<int *>(&current_state.algo),
algo_list, 2);
}

```

```

// testCnt++;

// if(testCnt == 5){
//     currentState.room_size = testCases[testCaseIdx];
//     testCaseIdx = (testCaseIdx+1) % 5;
//     testCnt = 0;
// }

if (currentState.algo == hdist::Algorithm::Sor) {
    ImGui::DragFloat("Sor Constant", &currentState.sor_constant, 0.01, 0.0, 20.0,
"%f");
}

// beginIteration = std::chrono::high_resolution_clock::now();

// broadcast current state
MPI_Bcast(&currentState, 1, MPI_State, 0, MPI_COMM_WORLD);

// change size regenerate
if (currentState.room_size != lastState.room_size) {
    grid = hdist::Grid{
        static_cast<size_t>(currentState.room_size),
        currentState.border_temp,
        currentState.source_temp,
        static_cast<size_t>(currentState.source_x),
        static_cast<size_t>(currentState.source_y)};
    first = true;
}

if (currentState != lastState) {
    lastState = currentState;
    finished = false;
}

if (first) {
    first = false;
    finished = false;
    begin = std::chrono::high_resolution_clock::now();
}

if (!finished) {
    int dataSize = currentState.room_size*currentState.room_size;

    // bcast get_current_buffer
    MPI_Bcast(grid.get_current_buffer().data(), dataSize, MPI_DOUBLE, 0,
MPI_COMM_WORLD);

    // do calculation
    finished = hdist::calculate(currentState, grid);
}

```

```

        bool localFinished = finished;

        // reduce and broadcast finished
        MPI_Reduce(&localFinished, &finished, 1, MPI_C_BOOL, MPI_LAND, 0,
MPI_COMM_WORLD);
        MPI_Bcast(&finished, 1, MPI_C_BOOL, 0, MPI_COMM_WORLD);

        // gatherv get_current_buffer result from all ranks
        // calculate the job size of each rank
        std::vector<int> taskSizeList;
        for(int i=0;i<processSize;i++){
            int startPosition = (dataSize/processSize)*i +
std::min(dataSize%processSize, i);
            int endPosition = (dataSize/processSize)*(i+1) +
std::min(dataSize%processSize, i+1)-1;
            taskSizeList.push_back(endPosition-startPosition+1);
        }

        // calculate the strips of received data packages
        std::vector<int> strips;
        strips.push_back(0);
        for(int i=1;i<processSize;i++){
            strips.push_back(strips[i-1]+taskSizeList[i-1]);
        }

        // calcualte job size of itself
        int startPosition = (dataSize/processSize)*rank +
std::min(dataSize%processSize, rank);
        int endPosition = (dataSize/processSize)*(rank+1) +
std::min(dataSize%processSize, rank+1)-1;
        int jobSize = endPosition-startPosition+1;

        // rank 0 gather calculated result
        MPI_Gatherv(grid.get_current_buffer().data()+startPosition, jobSize,
MPI_DOUBLE,
                grid.get_current_buffer().data(), taskSizeList.data(), strips.data(),
MPI_DOUBLE, 0, MPI_COMM_WORLD);

        if (finished) end = std::chrono::high_resolution_clock::now();

        // endIteration = std::chrono::high_resolution_clock::now();
        // printf("processes %d\ntime iteration %lld (ns)\ntroom size %d\n",
        // processSize,
std::chrono::duration_cast<std::chrono::nanoseconds>(endIteration - beginIteration).count(),
current_state.room_size);

    } else {

```

```

        ImGui::Text("stabilized in %lld ns",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
        printf("stabilized in %lld ns\n",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
    }

const ImVec2 p = ImGui::GetCursorScreenPos();
float x = p.x + current_state.block_size, y = p.y + current_state.block_size;
for (size_t i = 0; i < current_state.room_size; ++i) {
    for (size_t j = 0; j < current_state.room_size; ++j) {
        auto temp = grid[{i, j}];
        auto color = temp_to_color(temp);
        draw_list->AddRectFilled(ImVec2(x, y), ImVec2(x +
current_state.block_size, y + current_state.block_size), color);
        y += current_state.block_size;
    }
    x += current_state.block_size;
    y = p.y + current_state.block_size;
}
ImGui::End();
});

}

else{
    while(true){
        // receive current state
        MPI_Bcast(&current_state, 1, MPI_State, 0, MPI_COMM_WORLD);

        // change size regenerate
        if (current_state.room_size != last_state.room_size) {
            grid = hdist::Grid{
                static_cast<size_t>(current_state.room_size),
                current_state.border_temp,
                current_state.source_temp,
                static_cast<size_t>(current_state.source_x),
                static_cast<size_t>(current_state.source_y)};
            first = true;
        }

        if (current_state != last_state) {
            last_state = current_state;
            finished = false;
        }

        if (first) {
            first = false;
            finished = false;
        }

        if (!finished) {

```

```

        int dataSize = current_state.room_size*current_state.room_size;

        // bcast get_current_buffer
        MPI_Bcast(grid.get_current_buffer().data(), dataSize, MPI_DOUBLE, 0,
MPI_COMM_WORLD);

        // do calculation
        finished = hdist::calculate(current_state, grid);

        // reduce and receive finished
        MPI_Reduce(&finished, nullptr, 1, MPI_C_BOOL, MPI_LAND, 0, MPI_COMM_WORLD);
        MPI_Bcast(&finished, 1, MPI_C_BOOL, 0, MPI_COMM_WORLD);

        // send get_current_buffer result to rank 0
        // calcualte job size of itself
        int startPosition = (dataSize/processSize)*rank +
std::min(dataSize%processSize, rank);
        int endPosition = (dataSize/processSize)*(rank+1) +
std::min(dataSize%processSize, rank+1)-1;
        int jobSize = endPosition-startPosition+1;

        // send calculated result
        MPI_Gatherv(grid.get_current_buffer().data()+startPosition, jobSize,
MPI_DOUBLE, nullptr, nullptr, nullptr,
MPI_DOUBLE, 0, MPI_COMM_WORLD);
    }

}

// if plotting windows closed, rank 0 would then stop all other processes
if(rank==0){
    printf("aborting all processes\n");
    MPI_Abort(MPI_COMM_WORLD, 0);
}

MPI_Type_free(&MPI_State);
MPI_Finalize();
}

```

b) hdist.hpp

```

#pragma once

#include <vector>
#include <mpi.h>

namespace hdist {

```

```

enum class Algorithm : int {
    Jacobi = 0,
    Sor = 1
};

struct State {
    int room_size = 300;
    float block_size = 2;
    int source_x = room_size / 2;
    int source_y = room_size / 2;
    float source_temp = 100;
    float border_temp = 36;
    float tolerance = 0.02;
    float sor_constant = 4.0;
    Algorithm algo = hdist::Algorithm::Jacobi;

    bool operator==(const State &that) const = default;
};

struct Alt {

};

constexpr static inline Alt alt{};

struct Grid {
    std::vector<double> data0, data1;
    size_t current_buffer = 0;
    size_t length;

    explicit Grid(size_t size,
                  double border_temp,
                  double source_temp,
                  size_t x,
                  size_t y)
        : data0(size * size), data1(size * size), length(size) {
        for (size_t i = 0; i < length; ++i) {
            for (size_t j = 0; j < length; ++j) {
                if (i == 0 || j == 0 || i == length - 1 || j == length - 1) {
                    this->operator[](i, j) = border_temp;
                } else if (i == x && j == y) {
                    this->operator[](i, j) = source_temp;
                } else {
                    this->operator[](i, j) = 0;
                }
            }
        }
    }

    std::vector<double> &get_current_buffer() {

```

```

        if (current_buffer == 0) return data0;
        return data1;
    }

    std::vector<double> &get_alternate_buffer() {
        if (current_buffer == 0) return data1;
        return data0;
    }

    double &operator[](std::pair<size_t, size_t> index) {
        return get_current_buffer()[index.first * length + index.second];
    }

    double &operator[](std::tuple<Alt, size_t, size_t> index) {
        return current_buffer == 1 ? data0[std::get<1>(index) * length +
std::get<2>(index)] : data1[
            std::get<1>(index) * length + std::get<2>(index)];
    }

    void switch_buffer() {
        current_buffer = !current_buffer;
    }
};

struct UpdateResult {
    bool stable;
    double temp;
};

UpdateResult update_single(size_t i, size_t j, Grid &grid, const State &state) {
    UpdateResult result{};
    if (i == 0 || j == 0 || i == state.room_size - 1 || j == state.room_size - 1) {
        result.temp = state.border_temp;
    } else if (i == state.source_x && j == state.source_y) {
        result.temp = state.source_temp;
    } else {
        auto sum = (grid[{i + 1, j}] + grid[{i - 1, j}] + grid[{i, j + 1}] + grid[{i, j - 1}]);
        switch (state.algo) {
            case Algorithm::Jacobi:
                result.temp = 0.25 * sum;
                break;
            case Algorithm::Sor:
                result.temp = grid[{i, j}] + (1.0 / state.sor_constant) * (sum - 4.0 * grid[{i, j}]);
                break;
        }
        result.stable = fabs(grid[{i, j}] - result.temp) < state.tolerance;
    }
}

```

```

    return result;
}

bool calculate(const State &state, Grid &grid) {
    bool stabilized = true;
    int rank;
    int processSize;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &processSize);

    int dataSize = state.room_size*state.room_size;

    // calculate start and end position at buffer
    int startPosition = (dataSize/processSize)*rank + std::min(dataSize%processSize,
rank);
    int endPosition = (dataSize/processSize)*(rank+1) + std::min(dataSize%processSize,
rank+1)-1;
    int jobSize = endPosition-startPosition+1;

    // calculate start i and j
    int startI = startPosition/state.room_size;
    int i, j, count, endI;
    bool switchOnce = false;

    switch (state.algo) {
        case Algorithm::Jacobi:
            i = startI;
            j = startPosition%state.room_size;
            count = 0;
            for (; i < state.room_size; ++i) {
                if(i!=startI){
                    j=0;
                }
                for (; j < state.room_size; ++j) {
                    // avoid extra calculation
                    if(count>=jobSize){
                        // break outside for loop
                        i = state.room_size;
                        break;
                    }
                    count++;
                }
                auto result = update_single(i, j, grid, state);
                stabilized &= result.stable;
                grid[{alt, i, j}] = result.temp;
            }
        }
        grid.switch_buffer();
        break;
    }
}

```

```

case Algorithm::Sor:
    for (auto k : {0, 1}) {
        i = startI;
        j = startPosition%state.room_size;
        count = 0;

        // calculate the left and top extra points
        if(switchOnce==false){
            // left points
            if(i!=0){
                int x = i-1;
                for(int y=j;y<state.room_size;y++){
                    if (k == ((x + y) & 1)) {
                        auto result = update_single(x, y, grid, state);
                        // stabilized &= result.stable;
                        grid[{alt, x, y}] = result.temp;
                    } else {
                        grid[{alt, x, y}] = grid[{x, y}];
                    }
                }
            }

            // top points
            for(int y=0; y<j;y++){
                if (k == ((i + y) & 1)) {
                    auto result = update_single(i, y, grid, state);
                    // stabilized &= result.stable;
                    grid[{alt, i, y}] = result.temp;
                } else {
                    grid[{alt, i, y}] = grid[{i, y}];
                }
            }
        }

        bool breakLoop = false;

        for (; i < state.room_size; i++) {
            if(i!=startI){
                j=0;
            }

            for (; j < state.room_size; j++) {
                // avoid extra calculation
                if(count>=jobSize){
                    // break outside for loop
                    endI = i;
                    i = state.room_size;
                    break;
                }
            }
        }
    }
}

```

```

        }

        count++;

        if (k == ((i + j) & 1)) {
            auto result = update_single(i, j, grid, state);
            stabilized &= result.stable;
            grid[{alt, i, j}] = result.temp;
        } else {
            grid[{alt, i, j}] = grid[{i, j}];
        }
    }
}

// calculate the right and bottom extra points
if(switchOnce==false){
    // last point reach the bottom exactly
    int rightJ = j;
    int rightI = endI;
    if(j==0){
        rightJ = state.room_size-1;
        // block bottom points calculation
        j = state.room_size;
    }
    else{
        rightI++;
    }

    // calculate right points
    if(rightI<state.room_size){
        int x = rightI;

        for(int y=0;y<rightJ; y++){
            if (k == ((x + y) & 1)) {
                auto result = update_single(x, y, grid, state);
                // stabilized &= result.stable;
                grid[{alt, x, y}] = result.temp;
            } else {
                grid[{alt, x, y}] = grid[{x, y}];
            }
        }
    }

    // calculate bottom points
    for(int y=j;y<state.room_size;y++){
        if (k == ((endI + y) & 1)) {
            auto result = update_single(endI, y, grid, state);
            // stabilized &= result.stable;
            grid[{alt, endI, y}] = result.temp;
        } else {

```

```

        grid[{alt, endI, y}] = grid[{endI, y}];
    }
}

switchOnce = true;

grid.switch_buffer();
}
}

// printf("rank %d %d %d %d %d %d\n", rank, jobSize, starti, startj, curi, j);

return stabilized;
};

} // namespace hdist

```

2) Pthread

a) main.cpp

```

#include <graphic/graphic.hpp>
#include <imgui_impl_sdl.h>
#include <cstring>
#include <chrono>
#include <hdist/hdist.hpp>

template<typename ...Args>
void UNUSED(Args &&... args [[maybe_unused]]) {}

ImColor temp_to_color(double temp) {
    auto value = static_cast<uint8_t>(temp / 100.0 * 255.0);
    return {value, 0, 255 - value};
}

int main(int argc, char **argv) {
    // use 4 threads as default
    int totalThreadSize = 4;

    // fetch thread size from command line
    if(argc>=2){
        totalThreadSize = atoi(argv[1]);
    }

    // avoid invalid input
    if(totalThreadSize==0){

```

```

    totalThreadSize = 4;
}

printf("Pthread: launch %d threads\n", totalThreadSize);
// test codes
// int testCases[] = {200, 400, 800, 1600, 3200};
// int testCaseIdx = 0, testCnt = 0;
// static std::chrono::high_resolution_clock::time_point beginIteration, endIteration;

UNUSED(argc, argv);
bool first = true;
bool finished = false;
static hdist::State current_state, last_state;
static std::chrono::high_resolution_clock::time_point begin, end;
static const char* algo_list[2] = { "jacobi", "sor" };
graphic::GraphicContext context{"Assignment 4"};
auto grid = hdist::Grid{
    static_cast<size_t>(current_state.room_size),
    current_state.border_temp,
    current_state.source_temp,
    static_cast<size_t>(current_state.source_x),
    static_cast<size_t>(current_state.source_y)};
context.run([&](graphic::GraphicContext *context [[maybe_unused]], SDL_Window *) {
    auto io = ImGui::GetIO();
    ImGui::SetNextWindowPos(ImVec2(0.0f, 0.0f));
    ImGui::SetNextWindowSize(io.DisplaySize);
    ImGui::Begin("Assignment 4", nullptr,
        ImGuiWindowFlags_NoMove
        | ImGuiWindowFlags_NoCollapse
        | ImGuiWindowFlags_NoTitleBar
        | ImGuiWindowFlags_NoResize);
    ImDrawList *draw_list = ImGui::GetWindowDrawList();
    ImGui::Text("Application average %.3f ms/frame (%.1f FPS)", 1000.0f /
    ImGui::GetIO().Framerate,
        ImGui::GetIO().Framerate));
    ImGui::DragInt("Room Size", &current_state.room_size, 10, 200, 1600, "%d");
    ImGui::DragFloat("Block Size", &current_state.block_size, 0.01, 0.1, 10, "%f");
    ImGui::DragFloat("Source Temp", &current_state.source_temp, 0.1, 0, 100, "%f");
    ImGui::DragFloat("Border Temp", &current_state.border_temp, 0.1, 0, 100, "%f");
    ImGui::DragInt("Source X", &current_state.source_x, 1, 1, current_state.room_size - 2,
"%d");
    ImGui::DragInt("Source Y", &current_state.source_y, 1, 1, current_state.room_size - 2,
"%d");
    ImGui::DragFloat("Tolerance", &current_state.tolerance, 0.01, 0.01, 1, "%f");
    ImGui::ListBox("Algorithm", reinterpret_cast<int *>(&current_state.algo), algo_list,
2);

    // testCnt++;
}

```

```

// if(testCnt == 5){
//     currentState.room_size = testCases[testCaseIdx];
//     testCaseIdx = (testCaseIdx+1) % 5;
//     testCnt = 0;
// }

if (currentState.algo == hdist::Algorithm::Sor) {
    ImGui::DragFloat("Sor Constant", &currentState.sor_constant, 0.01, 0.0, 20.0,
"%f");
}

// beginIteration = std::chrono::high_resolution_clock::now();

if (currentState.room_size != lastState.room_size) {
    grid = hdist::Grid{
        static_cast<size_t>(currentState.room_size),
        currentState.border_temp,
        currentState.source_temp,
        static_cast<size_t>(currentState.source_x),
        static_cast<size_t>(currentState.source_y)};
    first = true;
}

if (currentState != lastState) {
    lastState = currentState;
    finished = false;
}

if (first) {
    first = false;
    finished = false;
    begin = std::chrono::high_resolution_clock::now();
}

if (!finished) {
    finished = hdist::calculate(currentState, grid, totalThreadSize);
    if (finished) end = std::chrono::high_resolution_clock::now();

    // endIteration = std::chrono::high_resolution_clock::now();
    // printf("processes %d\ntime iteration %lld (ns)\nroom size %d\n",
    // totalThreadSize,
std::chrono::duration_cast<std::chrono::nanoseconds>(endIteration - beginIteration).count(),
currentState.room_size);

} else {
    ImGui::Text("stabilized in %lld ns",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
    printf("stabilized in %lld ns\n",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
}

```

```

    }

    const ImVec2 p = ImGui::GetCursorScreenPos();
    float x = p.x + current_state.block_size, y = p.y + current_state.block_size;
    for (size_t i = 0; i < current_state.room_size; ++i) {
        for (size_t j = 0; j < current_state.room_size; ++j) {
            auto temp = grid[{i, j}];
            auto color = temp_to_color(temp);
            draw_list->AddRectFilled(ImVec2(x, y), ImVec2(x + current_state.block_size, y
+ current_state.block_size), color);
            y += current_state.block_size;
        }
        x += current_state.block_size;
        y = p.y + current_state.block_size;
    }
    ImGui::End();
});

}

```

b) hdist.hpp

```

#pragma once

#include <vector>
#include <iostream>

namespace hdist {
    // use to hold the barrier
    static pthread_barrier_t barrier;

    enum class Algorithm : int {
        Jacobi = 0,
        Sor = 1
    };

    struct State {
        int room_size = 300;
        float block_size = 2;
        int source_x = room_size / 2;
        int source_y = room_size / 2;
        float source_temp = 100;
        float border_temp = 36;
        float tolerance = 0.02;
        float sor_constant = 4.0;
        Algorithm algo = hdist::Algorithm::Jacobi;

        bool operator==(const State &that) const = default;
    };
}

```

```

struct Alt {
};

constexpr static inline Alt alt{};

struct Grid {
    std::vector<double> data0, data1;
    size_t current_buffer = 0;
    size_t length;

    explicit Grid(size_t size,
                  double border_temp,
                  double source_temp,
                  size_t x,
                  size_t y)
        : data0(size * size), data1(size * size), length(size) {
        for (size_t i = 0; i < length; ++i) {
            for (size_t j = 0; j < length; ++j) {
                if (i == 0 || j == 0 || i == length - 1 || j == length - 1) {
                    this->operator[](i, j) = border_temp;
                } else if (i == x && j == y) {
                    this->operator[](i, j) = source_temp;
                } else {
                    this->operator[](i, j) = 0;
                }
            }
        }
    }

    std::vector<double> &get_current_buffer() {
        if (current_buffer == 0) return data0;
        return data1;
    }

    double &operator[](std::pair<size_t, size_t> index) {
        return get_current_buffer()[index.first * length + index.second];
    }

    double &operator[](std::tuple<Alt, size_t, size_t> index) {
        return current_buffer == 1 ? data0[std::get<1>(index) * length +
std::get<2>(index)] : data1[
            std::get<1>(index) * length + std::get<2>(index)];
    }

    void switch_buffer() {
        current_buffer = !current_buffer;
    }
}

```

```

    static void switch_buffer(Grid & grid) {
        grid.current_buffer = !grid.current_buffer;
    }
};

// struct to hold parameters
struct Arguments{
    const State &state;
    Grid &grid;
    std::vector<int> &stateList;
    int taskNum;
    int totalThreadSize;
};

struct UpdateResult {
    bool stable;
    double temp;
};

UpdateResult update_single(size_t i, size_t j, Grid &grid, const State &state) {
    UpdateResult result{};
    if (i == 0 || j == 0 || i == state.room_size - 1 || j == state.room_size - 1) {
        result.temp = state.border_temp;
    } else if (i == state.source_x && j == state.source_y) {
        result.temp = state.source_temp;
    } else {
        auto sum = (grid[{i + 1, j}] + grid[{i - 1, j}] + grid[{i, j + 1}] + grid[{i, j - 1}]);
        switch (state.algo) {
            case Algorithm::Jacobi:
                result.temp = 0.25 * sum;
                break;
            case Algorithm::Sor:
                result.temp = grid[{i, j}] + (1.0 / state.sor_constant) * (sum - 4.0 * grid[{i, j}]);
                break;
        }
        result.stable = fabs(grid[{i, j}] - result.temp) < state.tolerance;
    }
    return result;
}

void* calculateInside(void *arg_ptr){
    // retrieve information from passed parameters
    auto arguments = static_cast<Arguments *>(arg_ptr);
    const State & state = arguments->state;
    Grid & grid = arguments->grid;
    std::vector<int> & stateList = arguments->stateList;
    int taskNum = arguments->taskNum;
}

```

```

int totalThreadSize = arguments->totalThreadSize;

// std::cout << &state << " " << &grid << " " << &stateList << "[" << taskNum << std::endl;

bool stabilized = true;

int dataSize = state.room_size*state.room_size;

// calculate start and end position at buffer
int startPosition = (dataSize/totalThreadSize)*taskNum +
std::min(dataSize%totalThreadSize, taskNum);
int endPosition = (dataSize/totalThreadSize)*(taskNum+1) +
std::min(dataSize%totalThreadSize, taskNum+1)-1;
int jobSize = endPosition-startPosition+1;

// calculate start i and j
int startI = startPosition/state.room_size;
int i, j, count;
bool switchOnce = false;

switch (state.algo) {
    case Algorithm::Jacobi:
        i = startI;
        j = startPosition%state.room_size;
        count = 0;
        for (; i < state.room_size; ++i) {
            if(i!=startI){
                j=0;
            }
            for (; j < state.room_size; ++j) {
                // avoid extra calculation
                if(count>=jobSize){
                    // break outside for loop
                    i = state.room_size;
                    break;
                }
                count++;
            }
            auto result = update_single(i, j, grid, state);
            stabilized &= result.stable;
            grid[{alt, i, j}] = result.temp;
        }
    }
    break;
    case Algorithm::Sor:
        for (auto k : {0, 1}) {
            i = startI;
            j = startPosition%state.room_size;
            count = 0;
        }
}

```

```

        for (; i < state.room_size; i++) {
            if(i!=startI){
                j=0;
            }
            for (; j < state.room_size; j++) {
                // avoid extra calculation
                if(count>=jobSize){
                    // break outside for loop
                    i = state.room_size;
                    break;
                }
                count++;
            }

            if (k == ((i + j) & 1)) {
                auto result = update_single(i, j, grid, state);
                stabilized &= result.stable;
                grid[{alt, i, j}] = result.temp;
            } else {
                grid[{alt, i, j}] = grid[{i, j}];
            }
        }
    }

    // wait until all threads finish calculation
    pthread_barrier_wait(&barrier);
    if(taskNum==0 && switchOnce==false){
        grid.switch_buffer();
        switchOnce = true;
    }
    pthread_barrier_wait(&barrier);
}

stateList[taskNum] = stabilized;

return nullptr;
}

bool calculate(const State &state, Grid &grid, int totalThreadSize) {
    std::vector<int> stateList(totalThreadSize, true);

    std::vector<pthread_t> threads(totalThreadSize);

    pthread_barrier_init(&barrier, NULL, totalThreadSize);

    // std::cout << &state << " " << &grid << " " << &stateList << std::endl;

    // initialize each thread with task information
    for(int i=0;i<totalThreadSize;i++){
        pthread_create(&threads[i], nullptr, calculateInside, new Arguments{

```

```

        .state = state,
        .grid = grid,
        .stateList = stateList,
        .taskNum = i,
        .totalThreadSize = totalThreadSize
    });
}

// wait until all the threads finish
for (auto & i : threads) {
    pthread_join(i, nullptr);
}

pthread_barrier_destroy(&barrier);

bool statelizedAll = true;

for(bool stabilized : stateList){
    // printf("checking %d\n", stabilized);
    statelizedAll &= stabilized;
}

grid.switch_buffer();

return statelizedAll;
}

} // namespace hdist

```

3) OpenMP:

a) main.cpp

```

#include <graphic/graphic.hpp>
#include <imgui_impl_sdl.h>
#include <cstring>
#include <chrono>
#include <hdist/hdist.hpp>

template<typename ...Args>
void UNUSED(Args &&... args [[maybe_unused]]) {}

ImColor temp_to_color(double temp) {
    auto value = static_cast<uint8_t>(temp / 100.0 * 255.0);
    return {value, 0, 255 - value};
}

```

```

int main(int argc, char **argv) {
    // use 4 threads as default
    int totalThreadSize = 4;

    // fetch thread size from command line
    if(argc>=2){
        totalThreadSize = atoi(argv[1]);
    }

    // avoid invalid input
    if(totalThreadSize==0){
        totalThreadSize = 4;
    }

    printf("OpenMP: launch %d threads\n", totalThreadSize);
    // test codes
    // int testCases[] = {200, 400, 800, 1600, 3200};
    // int testCaseIdx = 0, testCnt = 0;
    // static std::chrono::high_resolution_clock::time_point beginIteration, endIteration;

    omp_set_num_threads(totalThreadSize);

UNUSED(argc, argv);
bool first = true;
bool finished = false;
static hdist::State current_state, last_state;
static std::chrono::high_resolution_clock::time_point begin, end;
static const char* algo_list[2] = { "jacobi", "sor" };
graphic::GraphicContext context{"Assignment 4"};
auto grid = hdist::Grid{
    static_cast<size_t>(current_state.room_size),
    current_state.border_temp,
    current_state.source_temp,
    static_cast<size_t>(current_state.source_x),
    static_cast<size_t>(current_state.source_y)};
context.run([&](graphic::GraphicContext *context [[maybe_unused]], SDL_Window *) {
    auto io = ImGui::GetIO();
    ImGui::SetNextWindowPos(ImVec2(0.0f, 0.0f));
    ImGui::SetNextWindowSize(io.DisplaySize);
    ImGui::Begin("Assignment 4", nullptr,
                ImGuiWindowFlags_NoMove
                | ImGuiWindowFlags_NoCollapse
                | ImGuiWindowFlags_NoTitleBar
                | ImGuiWindowFlags_NoResize);
    ImDrawList *draw_list = ImGui::GetWindowDrawList();
    ImGui::Text("Application average %.3f ms/frame (%.1f FPS)", 1000.0f /
    ImGui::GetIO().Framerate,
               ImGui::GetIO().Framerate);
    ImGui::DragInt("Room Size", &current_state.room_size, 10, 200, 1600, "%d");
}

```

```

ImGui::DragFloat("Block Size", &current_state.block_size, 0.01, 0.1, 10, "%f");
ImGui::DragFloat("Source Temp", &current_state.source_temp, 0.1, 0, 100, "%f");
ImGui::DragFloat("Border Temp", &current_state.border_temp, 0.1, 0, 100, "%f");
ImGui::DragInt("Source X", &current_state.source_x, 1, 1, current_state.room_size - 2,
"%d");
ImGui::DragInt("Source Y", &current_state.source_y, 1, 1, current_state.room_size - 2,
"%d");
ImGui::DragFloat("Tolerance", &current_state.tolerance, 0.01, 0.01, 1, "%f");
ImGui::ListBox("Algorithm", reinterpret_cast<int *>(&current_state.algo), algo_list,
2);

// testCnt++;

// if(testCnt == 5){
//     current_state.room_size = testCases[testCaseIdx];
//     testCaseIdx = (testCaseIdx+1) % 5;
//     testCnt = 0;
// }

if (current_state.algo == hdist::Algorithm::Sor) {
    ImGui::DragFloat("Sor Constant", &current_state.sor_constant, 0.01, 0.0, 20.0,
"%f");
}

// beginIteration = std::chrono::high_resolution_clock::now();

if (current_state.room_size != last_state.room_size) {
    grid = hdist::Grid{
        static_cast<size_t>(current_state.room_size),
        current_state.border_temp,
        current_state.source_temp,
        static_cast<size_t>(current_state.source_x),
        static_cast<size_t>(current_state.source_y)};
    first = true;
}

if (current_state != last_state) {
    last_state = current_state;
    finished = false;
}

if (first) {
    first = false;
    finished = false;
    begin = std::chrono::high_resolution_clock::now();
}

if (!finished) {
    finished = hdist::calculate(current_state, grid, totalThreadSize);
}

```

```

    if (finished) end = std::chrono::high_resolution_clock::now();

    // endIteration = std::chrono::high_resolution_clock::now();
    // printf("thread %d\tone iteration %lld (ns)\troom size %d\n",
    // totalThreadSize,
std::chrono::duration_cast<std::chrono::nanoseconds>(endIteration - beginIteration).count(),
current_state.room_size);
} else {
    ImGui::Text("stabilized in %lld ns",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
    printf("stabilized in %lld ns\n",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
}

const ImVec2 p = ImGui::GetCursorScreenPos();
float x = p.x + current_state.block_size, y = p.y + current_state.block_size;
for (size_t i = 0; i < current_state.room_size; ++i) {
    for (size_t j = 0; j < current_state.room_size; ++j) {
        auto temp = grid[{i, j}];
        auto color = temp_to_color(temp);
        draw_list->AddRectFilled(ImVec2(x, y), ImVec2(x + current_state.block_size, y
+ current_state.block_size), color);
        y += current_state.block_size;
    }
    x += current_state.block_size;
    y = p.y + current_state.block_size;
}
ImGui::End();
});
}

```

b) hdist.hpp

```

#pragma once

#include <vector>
#include <omp.h>

namespace hdist {

enum class Algorithm : int {
    Jacobi = 0,
    Sor = 1
};

struct State {
    int room_size = 300;
    float block_size = 2;

```

```

int source_x = room_size / 2;
int source_y = room_size / 2;
float source_temp = 100;
float border_temp = 36;
float tolerance = 0.02;
float sor_constant = 4.0;
Algorithm algo = hdist::Algorithm::Jacobi;

bool operator==(const State &that) const = default;
};

struct Alt {
};

constexpr static inline Alt alt{};

struct Grid {
    std::vector<double> data0, data1;
    size_t current_buffer = 0;
    size_t length;

    explicit Grid(size_t size,
                  double border_temp,
                  double source_temp,
                  size_t x,
                  size_t y)
        : data0(size * size), data1(size * size), length(size) {
        for (size_t i = 0; i < length; ++i) {
            for (size_t j = 0; j < length; ++j) {
                if (i == 0 || j == 0 || i == length - 1 || j == length - 1) {
                    this->operator[]({i, j}) = border_temp;
                } else if (i == x && j == y) {
                    this->operator[]({i, j}) = source_temp;
                } else {
                    this->operator[]({i, j}) = 0;
                }
            }
        }
    }

    std::vector<double> &get_current_buffer() {
        if (current_buffer == 0) return data0;
        return data1;
    }

    double &operator[](std::pair<size_t, size_t> index) {
        return get_current_buffer()[index.first * length + index.second];
    }
}

```

```

        double &operator[](std::tuple<Alt, size_t, size_t> index) {
            return current_buffer == 1 ? data0[std::get<1>(index) * length +
std::get<2>(index)] : data1[
                std::get<1>(index) * length + std::get<2>(index)];
        }

        void switch_buffer() {
            current_buffer = !current_buffer;
        }
    };

    struct UpdateResult {
        bool stable;
        double temp;
    };
}

UpdateResult update_single(size_t i, size_t j, Grid &grid, const State &state) {
    UpdateResult result{};
    if (i == 0 || j == 0 || i == state.room_size - 1 || j == state.room_size - 1) {
        result.temp = state.border_temp;
    } else if (i == state.source_x && j == state.source_y) {
        result.temp = state.source_temp;
    } else {
        auto sum = (grid[{i + 1, j}] + grid[{i - 1, j}] + grid[{i, j + 1}] + grid[{i, j - 1}]);
        switch (state.algo) {
            case Algorithm::Jacobi:
                result.temp = 0.25 * sum;
                break;
            case Algorithm::Sor:
                result.temp = grid[{i, j}] + (1.0 / state.sor_constant) * (sum - 4.0 * grid[{i, j}]);
                break;
        }
        result.stable = fabs(grid[{i, j}] - result.temp) < state.tolerance;
    }
    return result;
}

bool calculate(const State &state, Grid &grid, int totalThreadSize) {
    std::vector<int> stateList(totalThreadSize, true);

    // total data size
    int dataSize = state.room_size * state.room_size;

    #pragma omp parallel
    {
        bool stabilized = true;

```

```

int taskNum = omp_get_thread_num();

// calculate start and end position at buffer
int startPosition = (dataSize/totalThreadSize)*taskNum +
std::min(dataSize%totalThreadSize, taskNum);
int endPosition = (dataSize/totalThreadSize)*(taskNum+1) +
std::min(dataSize%totalThreadSize, taskNum+1)-1;
int jobSize = endPosition-startPosition+1;

// calculate start i and j
int startI = startPosition/state.room_size;
int i, j, count;
bool switchOnce = false;

switch (state.algo) {
    case Algorithm::Jacobi:
        i = startI;
        j = startPosition%state.room_size;
        count = 0;
        for (; i < state.room_size; ++i) {
            if(i!=startI){
                j=0;
            }
            for (; j < state.room_size; ++j) {
                // avoid extra calculation
                if(count>=jobSize){
                    // break outside for loop
                    i = state.room_size;
                    break;
                }
                count++;
            }
            auto result = update_single(i, j, grid, state);
            stabilized &= result.stable;
            grid[{alt, i, j}] = result.temp;
        }
    }
    break;
    case Algorithm::Sor:
        for (auto k : {0, 1}) {
            i = startI;
            j = startPosition%state.room_size;
            count = 0;
            for (; i < state.room_size; i++) {
                if(i!=startI){
                    j=0;
                }
                for (; j < state.room_size; j++) {
                    // avoid extra calculation

```

```

        if(count>=jobSize){
            // break outside for loop
            i = state.room_size;
            break;
        }
        count++;

        if (k == ((i + j) & 1)) {
            auto result = update_single(i, j, grid, state);
            stabilized &= result.stable;
            grid[{alt, i, j}] = result.temp;
        } else {
            grid[{alt, i, j}] = grid[{i, j}];
        }
    }

    // wait until all threads finish calculation
#pragma omp barrier
if(taskNum==0 && switchOnce == false){
    grid.switch_buffer();
    switchOnce = true;
}
#pragma omp barrier
}

stateList[taskNum] = stabilized;

// printf("rank %d %d %d %d %d %d\n", taskNum, jobSize, starti, startj, curi,
j, stabilized);

}

bool statelizedAll = true;

for(bool stabilized : stateList){
    // printf("checking %d\n", stabilized);
    statelizedAll &= stabilized;
}

grid.switch_buffer();
return statelizedAll;
};

} // namespace hdist

```

4) CUDA

a) main.cu

```
#include <graphic/graphic.hpp>
#include <imgui_impl_sdl.h>
#include <cstring>
#include <chrono>
#include <hd़ist/hd़ist.hpp>

template<typename ...Args>
void UNUSED(Args &&... args [[maybe_unused]]) {}

ImColor temp_to_color(double temp) {
    auto value = static_cast<uint8_t>(temp / 100.0 * 255.0);
    return {value, 0, 255 - value};
}

int main(int argc, char **argv) {
    // use 4 threads as default
    int totalThreadSize = 4;

    // fetch thread size from command line
    if(argc>=2){
        totalThreadSize = atoi(argv[1]);
    }

    // avoid invalid input
    if(totalThreadSize==0){
        totalThreadSize = 4;
    }

    printf("CUDA: launch %d threads\n", totalThreadSize);
    // test codes
    // int testCases[] = {200, 400, 800, 1600, 3200};
    // int testCaseIdx = 0, testCnt = 0;
    // static std::chrono::high_resolution_clock::time_point beginIteration, endIteration;

    UNUSED(argc, argv);
    bool first = true;
    bool finished = false;
    static hd़ist::State current_state, last_state;
    static std::chrono::high_resolution_clock::time_point begin, end;
    static const char* algo_list[2] = { "jacobi", "sor" };
    graphic::GraphicContext context{"Assignment 4"};
    hd़ist::Grid * gridPtr = new hd़ist::Grid{
        static_cast<size_t>(current_state.room_size),
        current_state.border_temp,
        current_state.source_temp,
        static_cast<size_t>(current_state.source_x),
```

```

    static_cast<size_t>(current_state.source_y)};

hdist::Grid & grid = *gridPtr;
context.run([&](graphic::GraphicContext *context [[maybe_unused]], SDL_Window *) {
    auto io = ImGui::GetIO();
    ImGui::SetNextWindowPos(ImVec2(0.0f, 0.0f));
    ImGui::SetNextWindowSize(io.DisplaySize);
    ImGui::Begin("Assignment 4", nullptr,
        ImGuiWindowFlags_NoMove
        | ImGuiWindowFlags_NoCollapse
        | ImGuiWindowFlags_NoTitleBar
        | ImGuiWindowFlags_NoResize);
    ImDrawList *draw_list = ImGui::GetWindowDrawList();
    ImGui::Text("Application average %.3f ms/frame (%.1f FPS)", 1000.0f /
    ImGui::GetIO().Framerate,
        ImGui::GetIO().Framerate);
    ImGui::DragInt("Room Size", &current_state.room_size, 10, 200, 1600, "%d");
    ImGui::DragFloat("Block Size", &current_state.block_size, 0.01, 0.1, 10, "%f");
    ImGui::DragFloat("Source Temp", &current_state.source_temp, 0.1, 0, 100, "%f");
    ImGui::DragFloat("Border Temp", &current_state.border_temp, 0.1, 0, 100, "%f");
    ImGui::DragInt("Source X", &current_state.source_x, 1, 1, current_state.room_size - 2,
"%d");
    ImGui::DragInt("Source Y", &current_state.source_y, 1, 1, current_state.room_size - 2,
"%d");
    ImGui::DragFloat("Tolerance", &current_state.tolerance, 0.01, 0.01, 1, "%f");
    ImGui::ListBox("Algorithm", reinterpret_cast<int *>(&current_state.algo), algo_list,
2);

    // testCnt++;

    // if(testCnt == 5){
    //     current_state.room_size = testCases[testCaseIdx];
    //     testCaseIdx = (testCaseIdx+1) % 5;
    //     testCnt = 0;
    // }

    if (current_state.algo == hdist::Algorithm::Sor) {
        ImGui::DragFloat("Sor Constant", &current_state.sor_constant, 0.01, 0.0, 20.0,
"%f");
    }

    // beginIteration = std::chrono::high_resolution_clock::now();

    if (current_state.room_size != last_state.room_size) {
        delete gridPtr;
        gridPtr = new hdist::Grid{
            static_cast<size_t>(current_state.room_size),
            current_state.border_temp,
            current_state.source_temp,
            static_cast<size_t>(current_state.source_x),

```

```

        static_cast<size_t>(current_state.source_y));
grid = *gridPtr;
first = true;
}

if (!(current_state == last_state)) {
    last_state = current_state;
    finished = false;
}

if (first) {
    first = false;
    finished = false;
    begin = std::chrono::high_resolution_clock::now();
}

if (!finished) {
    finished = hdist::calculate(current_state, grid, totalThreadSize);
    if (finished) end = std::chrono::high_resolution_clock::now();

    // endIteration = std::chrono::high_resolution_clock::now();
    // printf("thread %d\ntone iteration %lld (ns)\troom size %d\n",
    // totalThreadSize,
std::chrono::duration_cast<std::chrono::nanoseconds>(endIteration - beginIteration).count(),
current_state.room_size);
} else {
    ImGui::Text("stabilized in %lld ns",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
    printf("stabilized in %lld ns\n",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
}

const ImVec2 p = ImGui::GetCursorScreenPos();
float x = p.x + current_state.block_size, y = p.y + current_state.block_size;
for (size_t i = 0; i < current_state.room_size; ++i) {
    for (size_t j = 0; j < current_state.room_size; ++j) {
        auto temp = grid.fetch(i, j);
        auto color = temp_to_color(temp);
        draw_list->AddRectFilled(ImVec2(x, y), ImVec2(x + current_state.block_size, y
+ current_state.block_size), color);
        y += current_state.block_size;
    }
    x += current_state.block_size;
    y = p.y + current_state.block_size;
}
ImGui::End();
});

}

```

b) hdist.hpp

```
#pragma once

namespace hdist {
    __device__
    int getBlockId() {
        return blockIdx.x + blockIdx.y * blockDim.x + blockDim.x * blockDim.y * blockIdx.z;
    }

    __device__
    int getLocalThreadId() {
        return (threadIdx.z * (blockDim.x * blockDim.y)) + (threadIdx.y * blockDim.x) +
    threadIdx.x;
    }

    __device__
    int getThreadId() {
        int blockId = getBlockId();
        int localThreadId = getLocalThreadId();
        return blockId * (blockDim.x * blockDim.y * blockDim.z) + localThreadId;
    }

    enum class Algorithm : int {
        Jacobi = 0,
        Sor = 1
    };

    struct State {
        int room_size = 300;
        float block_size = 2;
        int source_x = room_size / 2;
        int source_y = room_size / 2;
        float source_temp = 100;
        float border_temp = 36;
        float tolerance = 0.02;
        float sor_constant = 4.0;
        Algorithm algo = hdist::Algorithm::Jacobi;

        bool operator==(const State &that) const{
            return
                (this->room_size == that.room_size) &&
                (this->block_size == that.block_size) &&
                (this->source_x == that.source_x) &&
                (this->source_y == that.source_y) &&
                (this->source_temp == that.source_temp) &&
                (this->border_temp == that.border_temp) &&
                (this->tolerance == that.tolerance) &&
```

```

        (this->sor_constant == that.sor_constant) &&
        (this->algo == that.algo);
    }
};

struct Alt {
};

constexpr static inline Alt alt{};

struct Grid {
    double *data0, *data1;
    size_t current_buffer = 0;
    size_t length;

    __host__
    explicit Grid(size_t size,
                  double border_temp,
                  double source_temp,
                  size_t x,
                  size_t y)
        : length(size) {

        data0 = (double*)malloc(size*size*sizeof(double));
        data1 = (double*)malloc(size*size*sizeof(double));

        for (size_t i = 0; i < length; ++i) {
            for (size_t j = 0; j < length; ++j) {
                if (i == 0 || j == 0 || i == length - 1 || j == length - 1) {
                    this->set(i, j, border_temp);
                } else if (i == x && j == y) {
                    this->set(i, j, source_temp);
                } else {
                    this->set(i, j, 0);
                }
            }
        }
    }

    __host__ __device__
    double* get_current_buffer() {
        if (current_buffer == 0) return data0;
        return data1;
    }

    __host__ __device__
    double fetch(int index1, int index2) {
        return get_current_buffer()[index1 * length + index2];
    }
}

```

```

__host__ __device__
double fetch(Alt alt, int index1, int index2) {
    return current_buffer == 1 ? data0[index1 * length + index2] : data1[
        index1 * length + index2];
}

__host__ __device__
void set(int index1, int index2, double setVal) {
    get_current_buffer()[index1 * length + index2] = setVal;
}

__host__ __device__
void set(Alt alt, int index1, int index2, double setVal) {
    if(current_buffer == 1){
        data0[index1 * length + index2] = setVal;
    }
    else{
        data1[index1 * length + index2] = setVal;
    }
}

__host__ __device__
void switch_buffer() {
    current_buffer = !current_buffer;
}

__host__
~Grid(){
    free(data0);
    free(data1);
}
};

struct UpdateResult {
    bool stable;
    double temp;
};

__host__ __device__
UpdateResult update_single(size_t i, size_t j, Grid &grid, const State &state) {
    UpdateResult result{};
    if (i == 0 || j == 0 || i == state.room_size - 1 || j == state.room_size - 1) {
        result.temp = state.border_temp;
    } else if (i == state.source_x && j == state.source_y) {
        result.temp = state.source_temp;
    } else {
        auto sum = (grid.fetch(i + 1, j) + grid.fetch(i - 1, j) + grid.fetch(i, j + 1) +
grid.fetch(i, j - 1));
        result.temp = sum / 4;
    }
    result.stable = true;
}

```

```

switch (state.algo) {
    case Algorithm::Jacobi:
        result.temp = 0.25 * sum;
        break;
    case Algorithm::Sor:
        result.temp = grid.fetch(i, j) + (1.0 / state.sor_constant) * (sum - 4.0
* grid.fetch(i, j));
        break;
    }
    result.stable = fabs(grid.fetch(i, j) - result.temp) < state.tolerance;
    return result;
}

__global__
void cuda_calculate(const State &state, Grid &grid, int totalThreadSize, bool * stateList){
    // total data size
    int dataSize = state.room_size*state.room_size;

    bool stabilized = true;

    int taskNum = getThreadId();

    // calculate start and end position at buffer
    int startPosition = (dataSize/totalThreadSize)*taskNum + min(dataSize%totalThreadSize,
taskNum);
    int endPosition = (dataSize/totalThreadSize)*(taskNum+1) +
min(dataSize%totalThreadSize, taskNum+1)-1;
    int jobSize = endPosition-startPosition+1;

    // calculate start i and j
    int startI = startPosition/state.room_size;
    int i, j, count;
    bool switchOnce = false;

    switch (state.algo) {
        case Algorithm::Jacobi:
            i = startI;
            j = startPosition%state.room_size;
            count = 0;
            for (; i < state.room_size; ++i) {
                if(i!=startI){
                    j=0;
                }
                for (; j < state.room_size; ++j) {
                    // avoid extra calculation
                    if(count>=jobSize){
                        // break outside for loop
                        i = state.room_size;

```

```

        break;
    }
    count++;

    auto result = update_single(i, j, grid, state);
    stabilized &= result.stable;
    grid.set(alt, i, j, result.temp);
}
}

break;
case Algorithm::Sor:
for (auto k : {0, 1}) {
    i = startI;
    j = startPosition%state.room_size;
    count = 0;
    for (; i < state.room_size; i++) {
        if(i!=startI){
            j=0;
        }
        for (; j < state.room_size; j++) {
            // avoid extra calculation
            if(count>=jobSize){
                // break outside for loop
                i = state.room_size;
                break;
            }
            count++;
        }
        if (k == ((i + j) & 1)) {
            auto result = update_single(i, j, grid, state);
            stabilized &= result.stable;
            grid.set(alt, i, j, result.temp);
        } else {
            grid.set(alt, i, j, grid.fetch(i, j));
        }
    }
}

// make sure all threads finish calculation
__syncthreads();
if(taskNum==0 && switchOnce == false){
    grid.switch_buffer();
    switchOnce = true;
}
// make sure all threads finish calculation
__syncthreads();

}
}

```

```

        stateList[taskNum] = stabilized;
    }

__host__
bool calculate(const State &state, Grid &grid, int totalThreadSize) {

    // allocate state
    State * stateCopyPtr;
    cudaMallocManaged((void**)&stateCopyPtr, sizeof(state));
    cudaMemcpy(stateCopyPtr, &state, sizeof(state), cudaMemcpyHostToDevice);
    State & stateCopyReference = *stateCopyPtr;

    // allocate grid
    Grid * gridCopyPtr;
    double *data0, *data1;
    cudaMallocManaged((void**)&gridCopyPtr, sizeof(grid));
    cudaMallocManaged((void**)&data0, state.room_size*state.room_size*sizeof(double));
    cudaMallocManaged((void**)&data1, state.room_size*state.room_size*sizeof(double));
    cudaMemcpy(gridCopyPtr, &grid, sizeof(grid), cudaMemcpyHostToDevice);
    Grid & gridCopyReference = *gridCopyPtr;
    gridCopyReference.data0 = data0;
    gridCopyReference.data1 = data1;
    cudaMemcpy(gridCopyReference.get_current_buffer(), grid.get_current_buffer(),
state.room_size*state.room_size*sizeof(double), cudaMemcpyHostToDevice);

    // allocate stateList
    bool * stateList;
    cudaMallocManaged((void**)&stateList, totalThreadSize*sizeof(bool));
    memset(stateList, true, sizeof(stateList));

    cuda_calculate<<<1, totalThreadSize>>>(stateCopyReference, gridCopyReference,
totalThreadSize, stateList);
    cudaDeviceSynchronize();

    gridCopyReference.switch_buffer();
    grid.switch_buffer();

    // copy back the newest data
    memcpy(grid.get_current_buffer(), gridCopyReference.get_current_buffer(),
state.room_size*state.room_size*sizeof(double));

    bool statelizedAll = true;

    for(int i=0;i<totalThreadSize; i++){
        // printf("checking %d\n", stabilized);
        statelizedAll &= stateList[i];
    }
}

```

```

        // free allocated space
        cudaFree(stateCopyPtr);
        cudaFree(gridCopyPtr);
        cudaFree(data0);
        cudaFree(data1);
        cudaFree(stateList);

        return statelizedAll;
    };

} // namespace hdist

```

5) MPI with OpenMP

a) main.cpp

```

#include <graphic/graphic.hpp>
#include <imgui_impl_sdl.h>
#include <cstring>
#include <chrono>
#include <hdist/hdist.hpp>

template<typename ...Args>
void UNUSED(Args &&... args [[maybe_unused]]) {}

ImColor temp_to_color(double temp) {
    auto value = static_cast<uint8_t>(temp / 100.0 * 255.0);
    return {value, 0, 255 - value};
}

static MPI_Datatype MPI_State;

int main(int argc, char **argv) {
    // use 4 threads as default
    int totalThreadSize = 4;

    // fetch thread size from command line
    if(argc>=2){
        totalThreadSize = atoi(argv[1]);
    }

    // avoid invalid input
    if(totalThreadSize==0){
        totalThreadSize = 4;
    }

    int rank;
    int processSize;
    MPI_Init(&argc, &argv);
}

```

```

MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &processSize);

if(rank==0){
    printf("MPI+OpenMP: launch %d processes with %d threads\n", processSize,
totalThreadSize);
}

omp_set_num_threads(totalThreadSize);

// initialize MPI struct
const int nitems=9;
int blocklengths[9] = {1,1,1,1,1,1,1,1,1};
MPI_Datatype types[9] = {MPI_INT, MPI_FLOAT, MPI_INT, MPI_INT,
                        MPI_FLOAT, MPI_FLOAT, MPI_FLOAT, MPI_FLOAT, MPI_INT};

MPI_Aint offsets[9];

offsets[0] = offsetof(hdist::State, room_size);
offsets[1] = offsetof(hdist::State, block_size);
offsets[2] = offsetof(hdist::State, source_x);
offsets[3] = offsetof(hdist::State, source_y);
offsets[4] = offsetof(hdist::State, source_temp);
offsets[5] = offsetof(hdist::State, border_temp);
offsets[6] = offsetof(hdist::State, tolerance);
offsets[7] = offsetof(hdist::State, sor_constant);
offsets[8] = offsetof(hdist::State, algo);

MPI_Type_create_struct(nitems, blocklengths, offsets, types, &MPI_State);
MPI_Type_commit(&MPI_State);

UNUSED(argc, argv);
bool first = true;
bool finished = false;
static hdist::State current_state, last_state;
auto grid = hdist::Grid{
    static_cast<size_t>(current_state.room_size),
    current_state.border_temp,
    current_state.source_temp,
    static_cast<size_t>(current_state.source_x),
    static_cast<size_t>(current_state.source_y)};

if(rank == 0){
    // test codes
    // int testCases[] = {200, 400, 800, 1600, 3200};
    // int testCaseIdx = 0, testCnt = 0;
    // static std::chrono::high_resolution_clock::time_point beginIteration,
endIteration;
}

```

```

static std::chrono::high_resolution_clock::time_point begin, end;
static const char* algo_list[2] = { "jacobi", "sor" };
graphic::GraphicContext context{"Assignment 4"};
context.run([&](graphic::GraphicContext *context [[maybe_unused]], SDL_Window *) {
    auto io = ImGui::GetIO();
    ImGui::SetNextWindowPos(ImVec2(0.0f, 0.0f));
    ImGui::SetNextWindowSize(io.DisplaySize);
    ImGui::Begin("Assignment 4", nullptr,
                ImGuiWindowFlags_NoMove
                | ImGuiWindowFlags_NoCollapse
                | ImGuiWindowFlags_NoTitleBar
                | ImGuiWindowFlags_NoResize);
    ImDrawList *draw_list = ImGui::GetWindowDrawList();
    ImGui::Text("Application average %.3f ms/frame (%.1f FPS)", 1000.0f /
    ImGui::GetIO().Framerate,
               ImGui::GetIO().Framerate);
    ImGui::DragInt("Room Size", &current_state.room_size, 10, 200, 1600, "%d");
    ImGui::DragFloat("Block Size", &current_state.block_size, 0.01, 0.1, 10, "%f");
    ImGui::DragFloat("Source Temp", &current_state.source_temp, 0.1, 0, 100, "%f");
    ImGui::DragFloat("Border Temp", &current_state.border_temp, 0.1, 0, 100, "%f");
    ImGui::DragInt("Source X", &current_state.source_x, 1, 1, current_state.room_size
- 2, "%d");
    ImGui::DragInt("Source Y", &current_state.source_y, 1, 1, current_state.room_size
- 2, "%d");
    ImGui::DragFloat("Tolerance", &current_state.tolerance, 0.01, 0.01, 1, "%f");
    ImGui::ListBox("Algorithm", reinterpret_cast<int *>(&current_state.algo),
algo_list, 2);

    // testCnt++;

    // if(testCnt == 5){
    //     current_state.room_size = testCases[testCaseIdx];
    //     testCaseIdx = (testCaseIdx+1) % 5;
    //     testCnt = 0;
    // }

    if (current_state.algo == hdist::Algorithm::Sor) {
        ImGui::DragFloat("Sor Constant", &current_state.sor_constant, 0.01, 0.0, 20.0,
"%f");
    }

    // beginIteration = std::chrono::high_resolution_clock::now();

    // broadcast current state
    MPI_Bcast(&current_state, 1, MPI_State, 0, MPI_COMM_WORLD);

    // change size regenerate
    if (current_state.room_size != last_state.room_size) {
        grid = hdist::Grid{

```

```

        static_cast<size_t>(current_state.room_size),
        current_state.border_temp,
        current_state.source_temp,
        static_cast<size_t>(current_state.source_x),
        static_cast<size_t>(current_state.source_y)};
    first = true;
}

if (current_state != last_state) {
    last_state = current_state;
    finished = false;
}

if (first) {
    first = false;
    finished = false;
    begin = std::chrono::high_resolution_clock::now();
}

if (!finished) {
    int dataSize = current_state.room_size*current_state.room_size;

    // bcast get_current_buffer
    MPI_Bcast(grid.get_current_buffer().data(), dataSize, MPI_DOUBLE, 0,
MPI_COMM_WORLD);

    // do calculation
    finished = hdist::calculate(current_state, grid, totalThreadSize);

    bool localFinished = finished;

    // reduce and broadcast finished
    MPI_Reduce(&localFinished, &finished, 1, MPI_C_BOOL, MPI_LAND, 0,
MPI_COMM_WORLD);
    MPI_Bcast(&finished, 1, MPI_C_BOOL, 0, MPI_COMM_WORLD);

    // gatherv get_current_buffer result from all ranks
    // calculate the job size of each rank
    std::vector<int> taskSizeList;
    for(int i=0;i<processSize;i++){
        int startPosition = (dataSize/processSize)*i +
std::min(dataSize%processSize, i);
        int endPosition = (dataSize/processSize)*(i+1) +
std::min(dataSize%processSize, i+1)-1;
        taskSizeList.push_back(endPosition-startPosition+1);
    }

    // calculate the strips of received data packages
    std::vector<int> strips;
}

```

```

        strips.push_back(0);
        for(int i=1;i<processSize;i++){
            strips.push_back(strips[i-1]+taskSizeList[i-1]);
        }

        // calcualte job size of itself
        int startPosition = (dataSize/processSize)*rank +
std::min(dataSize%processSize, rank);
        int endPosition = (dataSize/processSize)*(rank+1) +
std::min(dataSize%processSize, rank+1)-1;
        int jobSize = endPosition-startPosition+1;

        // rank 0 gather calculated result
        MPI_Gatherv(grid.get_current_buffer().data()+startPosition, jobSize,
MPI_DOUBLE,
            grid.get_current_buffer().data(), taskSizeList.data(), strips.data(),
MPI_DOUBLE, 0, MPI_COMM_WORLD);

        if (finished) end = std::chrono::high_resolution_clock::now();

        // endIteration = std::chrono::high_resolution_clock::now();
        // printf("threads %d\ntime iteration %lld (ns)\nroom size %d\n",
        // processSize*totalThreadSize,
std::chrono::duration_cast<std::chrono::nanoseconds>(endIteration - beginIteration).count(),
current_state.room_size);

    } else {
        ImGui::Text("stabilized in %lld ns",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
        printf("stabilized in %lld ns\n",
std::chrono::duration_cast<std::chrono::nanoseconds>(end - begin).count());
    }

    const ImVec2 p = ImGui::GetCursorScreenPos();
    float x = p.x + current_state.block_size, y = p.y + current_state.block_size;
    for (size_t i = 0; i < current_state.room_size; ++i) {
        for (size_t j = 0; j < current_state.room_size; ++j) {
            auto temp = grid[{i, j}];
            auto color = temp_to_color(temp);
            draw_list->AddRectFilled(ImVec2(x, y), ImVec2(x +
current_state.block_size, y + current_state.block_size), color);
            y += current_state.block_size;
        }
        x += current_state.block_size;
        y = p.y + current_state.block_size;
    }
    ImGui::End();
});
}

```

```

else{
    while(true){
        // receive current state
        MPI_Bcast(&current_state, 1, MPI_State, 0, MPI_COMM_WORLD);

        // change size regenerate
        if (current_state.room_size != last_state.room_size) {
            grid = hdist::Grid{
                static_cast<size_t>(current_state.room_size),
                current_state.border_temp,
                current_state.source_temp,
                static_cast<size_t>(current_state.source_x),
                static_cast<size_t>(current_state.source_y)};
            first = true;
        }

        if (current_state != last_state) {
            last_state = current_state;
            finished = false;
        }

        if (first) {
            first = false;
            finished = false;
        }

        if (!finished) {
            int dataSize = current_state.room_size*current_state.room_size;

            // bcast get_current_buffer
            MPI_Bcast(grid.get_current_buffer().data(), dataSize, MPI_DOUBLE, 0,
MPI_COMM_WORLD);

            // do calculation
            finished = hdist::calculate(current_state, grid, totalThreadSize);

            // reduce and receive finished
            MPI_Reduce(&finished, nullptr, 1, MPI_C_BOOL, MPI_LAND, 0, MPI_COMM_WORLD);
            MPI_Bcast(&finished, 1, MPI_C_BOOL, 0, MPI_COMM_WORLD);

            // send get_current_buffer result to rank 0
            // calcualte job size of itself
            int startPosition = (dataSize/processSize)*rank +
std::min(dataSize%processSize, rank);
            int endPosition = (dataSize/processSize)*(rank+1) +
std::min(dataSize%processSize, rank+1)-1;
            int jobSize = endPosition-startPosition+1;

            // send calculated result
        }
    }
}

```

```

        MPI_Gatherv(grid.get_current_buffer().data()+startPosition, jobSize,
MPI_DOUBLE, nullptr, nullptr, nullptr,
                MPI_DOUBLE, 0, MPI_COMM_WORLD);
}

}

// if plotting windows closed, rank 0 would then stop all other processes
if(rank==0){
    printf("aborting all processes\n");
    MPI_Abort(MPI_COMM_WORLD, 0);
}

MPI_Type_free(&MPI_State);
MPI_Finalize();
}

```

b) hdist.hpp

```

#pragma once

#include <vector>
#include <mpi.h>
#include <omp.h>

namespace hdist {

enum class Algorithm : int {
    Jacobi = 0,
    Sor = 1
};

struct State {
    int room_size = 300;
    float block_size = 2;
    int source_x = room_size / 2;
    int source_y = room_size / 2;
    float source_temp = 100;
    float border_temp = 36;
    float tolerance = 0.02;
    float sor_constant = 4.0;
    Algorithm algo = hdist::Algorithm::Jacobi;

    bool operator==(const State &that) const = default;
};

struct Alt {

```

```
};

constexpr static inline Alt alt{};

struct Grid {
    std::vector<double> data0, data1;
    size_t current_buffer = 0;
    size_t length;

    explicit Grid(size_t size,
                  double border_temp,
                  double source_temp,
                  size_t x,
                  size_t y)
        : data0(size * size), data1(size * size), length(size) {
        for (size_t i = 0; i < length; ++i) {
            for (size_t j = 0; j < length; ++j) {
                if (i == 0 || j == 0 || i == length - 1 || j == length - 1) {
                    this->operator[](i, j) = border_temp;
                } else if (i == x && j == y) {
                    this->operator[](i, j) = source_temp;
                } else {
                    this->operator[](i, j) = 0;
                }
            }
        }
    }

    std::vector<double> &get_current_buffer() {
        if (current_buffer == 0) return data0;
        return data1;
    }

    std::vector<double> &get_alternate_buffer() {
        if (current_buffer == 0) return data1;
        return data0;
    }

    double &operator[](std::pair<size_t, size_t> index) {
        return get_current_buffer()[index.first * length + index.second];
    }

    double &operator[](std::tuple<Alt, size_t, size_t> index) {
        return current_buffer == 1 ? data0[std::get<1>(index) * length +
std::get<2>(index)] : data1[
            std::get<1>(index) * length + std::get<2>(index)];
    }

    void switch_buffer() {
```

```

        current_buffer = !current_buffer;
    }
};

struct UpdateResult {
    bool stable;
    double temp;
};

UpdateResult update_single(size_t i, size_t j, Grid &grid, const State &state) {
    UpdateResult result{};
    if (i == 0 || j == 0 || i == state.room_size - 1 || j == state.room_size - 1) {
        result.temp = state.border_temp;
    } else if (i == state.source_x && j == state.source_y) {
        result.temp = state.source_temp;
    } else {
        auto sum = (grid[{i + 1, j}] + grid[{i - 1, j}] + grid[{i, j + 1}] + grid[{i, j - 1}]);
        switch (state.algo) {
            case Algorithm::Jacobi:
                result.temp = 0.25 * sum;
                break;
            case Algorithm::Sor:
                result.temp = grid[{i, j}] + (1.0 / state.sor_constant) * (sum - 4.0 * grid[{i, j}]);
                break;
        }
    }
    result.stable = fabs(grid[{i, j}] - result.temp) < state.tolerance;
    return result;
}

bool calculate(const State &state, Grid &grid, int totalThreadSize) {
    bool stabilized = true;
    int rank;
    int processSize;
    std::vector<int> stateList(totalThreadSize, true);

    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &processSize);

    int dataSize = state.room_size*state.room_size;

    // calculate start and end position at buffer
    int processStartPosition = (dataSize/processSize)*rank +
std::min(dataSize%processSize, rank);
    int processEndPosition = (dataSize/processSize)*(rank+1) +
std::min(dataSize%processSize, rank+1)-1;
    int processTaskSize = processEndPosition - processStartPosition + 1;
}

```

```

#pragma omp parallel
{
    int taskNum = omp_get_thread_num();

    int startPosition = (processTaskSize/totalThreadSize)*taskNum +
std::min(processTaskSize%totalThreadSize, taskNum) + processStartPosition;
    int endPosition = (processTaskSize/totalThreadSize)*(taskNum+1) +
std::min(processTaskSize%totalThreadSize, taskNum+1)-1 + processStartPosition;
    int jobSize = endPosition-startPosition+1;

    // calculate start i and j
    int startI = startPosition/state.room_size;
    int i, j, count, endI;
    bool switchOnce = false;
    int starti, startj, curi;

    switch (state.algo) {
        case Algorithm::Jacobi:
            i = startI;
            j = startPosition%state.room_size;
            starti = i;
            startj = j;
            count = 0;
            for (; i < state.room_size; ++i) {
                if(i!=startI){
                    j=0;
                }
                curi = i;
                for (; j < state.room_size; ++j) {
                    // avoid extra calculation
                    if(count>=jobSize){
                        // break outside for loop
                        i = state.room_size;
                        break;
                    }
                    count++;
                }

                auto result = update_single(i, j, grid, state);
                stabilized &= result.stable;
                grid[{alt, i, j}] = result.temp;
            }
        }
        break;
    case Algorithm::Sor:
        for (auto k : {0, 1}) {
            i = startI;
            j = startPosition%state.room_size;
            count = 0;
        }
    }
}

```

```

// calculate the left and top extra points
if(switchOnce==false){
    // left points
    if(i!=0){
        int x = i-1;
        for(int y=j;y<state.room_size;y++){
            if (k == ((x + y) & 1)) {
                auto result = update_single(x, y, grid, state);
                // stabilized &= result.stable;
                grid[{alt, x, y}] = result.temp;
            } else {
                grid[{alt, x, y}] = grid[{x, y}];
            }
        }
    }

    // top points
    for(int y=0; y<j;y++){
        if (k == ((i + y) & 1)) {
            auto result = update_single(i, y, grid, state);
            // stabilized &= result.stable;
            grid[{alt, i, y}] = result.temp;
        } else {
            grid[{alt, i, y}] = grid[{i, y}];
        }
    }
}

for (; i < state.room_size; i++) {
    if(i!=startI){
        j=0;
    }

    for (; j < state.room_size; j++) {
        // avoid extra calculation
        if(count>=jobSize){
            // break outside for loop
            endI = i;
            i = state.room_size;
            break;
        }
        count++;

        if (k == ((i + j) & 1)) {
            auto result = update_single(i, j, grid, state);
            stabilized &= result.stable;
            grid[{alt, i, j}] = result.temp;
        } else {
    }
}

```

```

        grid[{alt, i, j}] = grid[{i, j}];
    }
}

// calculate the right and bottom extra points
if(switchOnce==false){
    // last point reach the bottom exactly
    int rightJ = j;
    int rightI = endI;
    if(j==0){
        rightJ = state.room_size-1;
        // block bottom points calculation
        j = state.room_size;
    }
    else{
        rightI++;
    }

    // calculate right points
    if(rightI<state.room_size){
        int x = rightI;

        for(int y=0;y<rightJ; y++){
            if (k == ((x + y) & 1)) {
                auto result = update_single(x, y, grid, state);
                // stabilized &= result.stable;
                grid[{alt, x, y}] = result.temp;
            } else {
                grid[{alt, x, y}] = grid[{x, y}];
            }
        }
    }

    // calculate bottom points
    for(int y=j;y<state.room_size;y++){
        if (k == ((endI + y) & 1)) {
            auto result = update_single(endI, y, grid, state);
            // stabilized &= result.stable;
            grid[{alt, endI, y}] = result.temp;
        } else {
            grid[{alt, endI, y}] = grid[{endI, y}];
        }
    }
}

// wait until all threads finish calculation
#pragma omp barrier

```

```
    if(taskNum==0 && switchOnce == false){
        grid.switch_buffer();
    }

    switchOnce = true;

    #pragma omp barrier

}

}

stateList[taskNum] = stabilized;

// printf("%d %d %d %d %d [%d]\n", rank, taskNum, starti, startj, curi, j,
stabilized);

}

bool statelizedAll = true;

for(bool stabilized : stateList){
    // printf("checking %d\n", stabilized);
    statelizedAll &= stabilized;
}

grid.switch_buffer();

return statelizedAll;
};

} // namespace hdist
```